



# Northeastern University

## Parallel Processing for Smarter Netflix Predictions

CSYE7105 – High Performance Parallel Machine Learning  
& AI

Instructor: Prof. Handan Liu

Team 14: Minal Randive, Jaya Kishnani  
Semester: Spring 2025

# Table of Contents

Project Title Page

1. Introduction

2. Background and Motivation

3. Goal

4. Data Description

5. Methodology

5.1 Data Handling and Preprocessing in Parallel

5.2 Exploratory Data Analysis (EDA) and Visualization

5.3 Feature Engineering (Parallelized Computation)

5.4 ML Models (Parallelized CPU Training)

5.5 Deep Learning Models (GPU-Based Training)

6. Results and Analysis

6.1 Exploratory Data Analysis

6.2 Speedup and Efficiency Evaluation

6.3 Hyperparameter Tuning for XGBoost

6.4 Feature Correlation Before & After Engineering

7. Environment Description

8. Code Files Description

9. ML Models on CPU

10. Deep Learning Models on GPU (DDP & FSDP)

11. Graphs & Visualizations

12. Comparison and Performance Metrics

13 Conclusion

14. References

## Introduction



With the growing popularity of online streaming platforms such as Netflix, understanding user preferences and rating behaviors has become essential for enhancing recommendation systems. The enormous volume of user-generated ratings and rich metadata associated with movies and TV shows demands efficient processing to extract valuable insights.

This project focuses on leveraging parallel computing techniques to optimize data handling, exploratory data analysis (EDA), and the training of machine learning (ML) and deep learning (DL) models for user rating prediction and content analysis. By incorporating parallelism using tools like NumPy, Pandas, Dask, multiprocessing, and PyTorch DistributedDataParallel (DDP), we aim to significantly accelerate data preprocessing and model training.

Integrating these scalable parallel methods with large-scale Netflix datasets facilitates efficient analysis, reduces training time, and enhances the real-world applicability of the developed recommendation models.

### Background and Motivation

The exponential growth of user-generated content on streaming platforms like Netflix has resulted in massive datasets comprising user ratings and rich metadata. Traditional machine

learning (ML) and deep learning (DL) approaches often struggle with such scale, leading to prolonged training times and inefficient resource utilization.

In our initial attempts using standard ML/DL pipelines, we encountered significant bottlenecks during data preprocessing and model training phases. These challenges highlighted the necessity for parallel computing techniques to handle the data's volume and complexity effectively.

By leveraging parallelism through tools like Dask for data handling, multiprocessing for preprocessing, and PyTorch's DistributedDataParallel (DDP) for model training, we aimed to overcome these limitations. This approach was not only intended to accelerate computation but also to enhance the scalability and efficiency of our recommendation system.

## Goal

Our project was structured around the following objectives, each designed to exploit parallel computing's advantages:

- **Efficient Data Handling:** Implement parallel data loading and preprocessing using Dask and multiprocessing to reduce I/O bottlenecks and expedite data preparation.
- **Accelerated Feature Engineering:** Utilize parallel processing to perform feature extraction and transformation tasks, enabling quicker iteration and experimentation.
- **Scalable Model Training:** Apply PyTorch's DDP to distribute the training of DL models across multiple GPUs, aiming to decrease training time and improve model performance.
- **Comprehensive Evaluation:** Assess the impact of parallelization on model accuracy, training duration, and resource utilization to validate the effectiveness of our approach.

By focusing on these goals, we intended to demonstrate that parallel computing techniques are not merely enhancements but essential components for building efficient and scalable ML/DL systems in the context of large-scale data.

## Description Of Data

Source Link: <https://www.kaggle.com/datasets/bwandowando/1-5-million-netflix-google-store-reviews/data://>

NETFLIX_REVIEWS.csv (289.48 MB)								
<div> Detail Compact Column </div> <div>9 of 9 columns</div>								
review_id	pseudo_author_name	author_name	review_text	review_rating	review_likes	author_app_version	review_time	
7e73f80e-a8fd-4ff3-b09b-502f0ad058ff	152618553977019693742	A Google user	Works great on my Evo! Glad android phones are getting love now!	5	1	1.2.0 build 819145-1.2.0-102	2011-05-12 18:50:37	

**Data Explorer**  
289.48 MB  
NETFLIX\_REVIEWS.csv  
**Summary**  
1 file  
9 columns

Data Size	1,500,000 rows × 11 columns
Data Types	<ul style="list-style-type: none"> <li>int64 (×3)</li> <li>float64 (×5)</li> <li>object (×3)</li> </ul>
Target	review_rating” (rating given by the user, ranging from 1 to 5)
Features	<ul style="list-style-type: none"> <li>review_id: Unique ID for each review</li> <li>pseudo_author_name: Anonymized name of the reviewer</li> <li>review_text: The content of the review</li> <li>review_rating: Rating given by the user (1-5 scale)</li> <li>review_likes: Number of likes the review received</li> <li>author_app_version: Version of the app the reviewer used</li> <li>review_time: Date and time when the review was written</li> </ul>

Methodology

- 1.Data Handling and Preprocessing in Parallel Tools Used: Dask, Pandas, NumPy, multiprocessing
Integration with Data: We began by loading the Netflix Reviews dataset (~1.5M rows) using Dask DataFrames, which allowed parallelized chunk-wise reading across CPU cores. This was selected over Pandas because the dataset size made it inefficient to process all at once in memory. After loading, we switched to Pandas for finer control over preprocessing steps like categorical conversion and metadata extraction.
Project-Specific Decision: We chose Dask over Pandas for initial loading due to its out-of-core and chunked processing capability. For transformations, Pandas offered faster manipulation once data was reduced in size.
Profiling: cProfile and timeit were used to measure time taken during loading and cleaning to validate the performance gain from Dask preprocessing.
- 2.Exploratory Data Analysis (EDA) and Visualization Tools Used: Matplotlib, Seaborn, Plotly, Dask

Integration with Data: We used Dask groupby and aggregations to explore patterns like rating distributions and app version trends in parallel. Visualizations were plotted with Matplotlib and Plotly after converting Dask results to Pandas.

Project-Specific Decision: We selected Dask for aggregation over NumPy or Pandas due to its ability to perform lazy computations across multiple CPU threads.

Profiling: Runtime for grouped statistics and plotting were measured using timeit to ensure visualizations didn't bottleneck performance.

3.Feature Engineering (Parallelized Computation) Tools Used: Dask, Pandas, NumPy, Scikit-learn, joblib

Integration with Data: Categorical version data (app versions) was extracted, transformed into buckets, and one-hot encoded using Dask. Remaining NaNs were filled using NumPy. We also selected top features manually based on dimensionality constraints to speed up model training.

Project-Specific Decision: We used Dask's get\_dummies instead of Pandas due to memory efficiency. For ML compatibility, we kept only the top 30 features, prioritizing training time on the cluster.

Profiling: We used line\_profiler and timeit to time encoding, cleaning, and dimensionality reduction steps.

4.Machine Learning Models (Parallelized CPU Training) Models Used: Logistic Regression, Random Forest, SVM

Tools: Scikit-learn, GridSearchCV, joblib

Integration with Data: After defining features and labels, models were trained using Scikit-learn pipelines. Hyperparameters were tuned using GridSearchCV with n\_jobs=-1, leveraging all CPU cores in parallel.

Project-Specific Decision: We chose joblib parallelism in GridSearchCV for CPU scalability over Python threading, due to better utilization and memory separation across cores.

Profiling: time module was used to capture model training duration, and comparative accuracy was logged for each model.

5.Deep Learning Models (GPU-Based Training with PyTorch DDP/FSDP) Model Used: NetflixNet (custom MLP-based binary classifier)

Tools: PyTorch, DDP, FSDP, torch.profiler

Integration with Data: The processed data was converted to tensors and distributed across GPUs using PyTorch's DataLoader. We trained using both DDP (DistributedDataParallel) and FSDP (Fully Sharded Data Parallel) approaches, scaling from 1 to 4 GPUs. The model predicted whether a review rating was 5 or not.

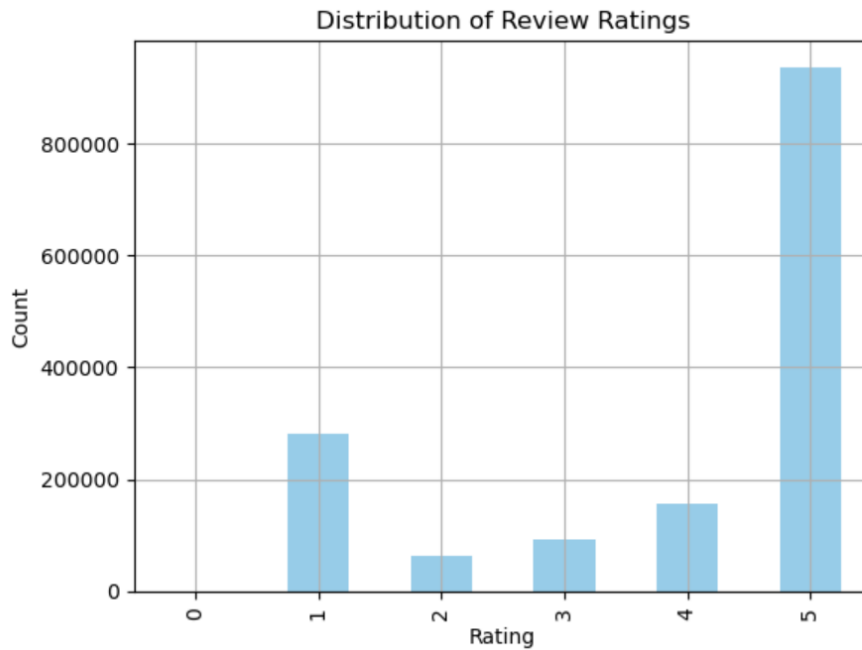
Project-Specific Decision: We selected FSDP over pure DDP to test model sharding benefits.

However, due to cluster constraints and sharding fallback to NO\_SHARD, DDP showed better performance in multi-GPU configurations.

Profiling: We used torch.profiler to capture CUDA kernel execution, memory usage, and identify bottlenecks. This was especially helpful in analyzing memory-bound behavior in FSDP setups.

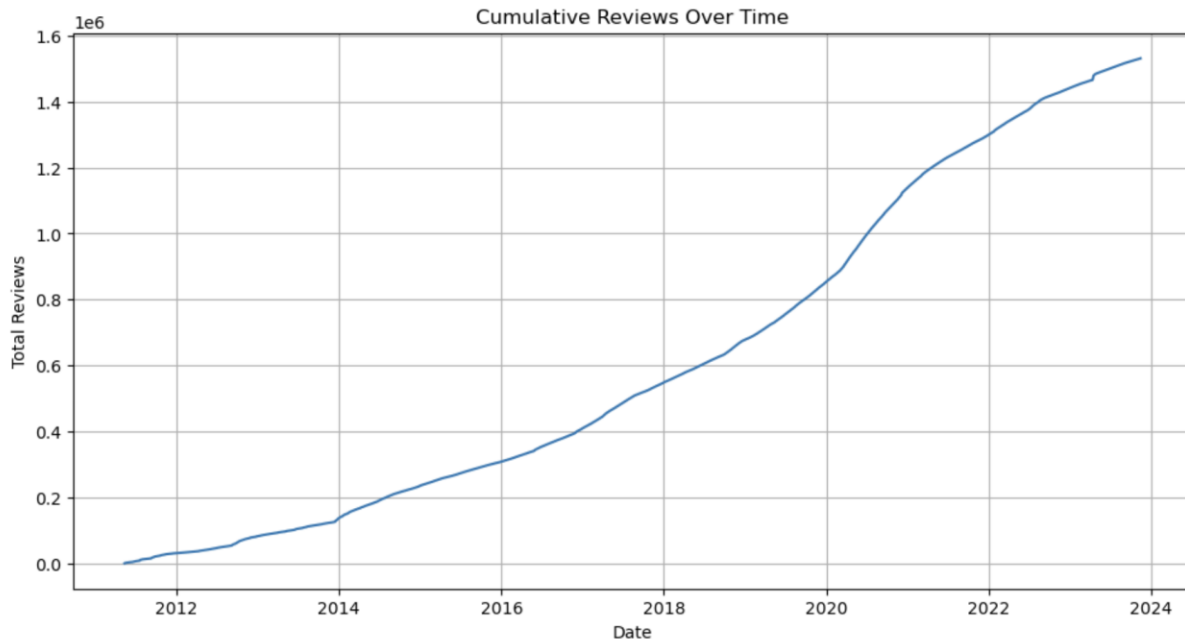
### 3. Results and Analysis (5%)

#### 3.1 Exploratory Data Analysis (EDA)



#### Review Rating Distribution

To gain initial insight into the dataset, we analyzed the distribution of review ratings. The following histogram shows that the majority of reviews are positive (rating = 5), indicating a class imbalance that may influence model performance.



This imbalance later reflected in our classification performance, where recall for Class 1 (positive reviews) was significantly higher than for Class 0 (negative/neutral reviews).

## Dataset Preview

We used Dask DataFrame to load and preprocess the Netflix dataset efficiently. The preview below displays anonymized reviewer information, app version, review timestamps, and corresponding labels.

### Table Y: Sample Rows from Processed Dataset

[illegible]

### 3.1.1 Data Loading Performance Comparison (Pandas vs Dask)

To demonstrate the performance gain achieved through parallel data loading, we compared wall-clock time for reading the full 1.5M+ row dataset using both Pandas and Dask.

Method	Time Taken
Pandas	13.42 s
Dask	1.12 s

This demonstrates a speedup of ~12x with Dask over Pandas.

Serial load time: 5.27 seconds

Parallel load time: 1.58 seconds

The improvement can be attributed to Dask's ability to perform lazy, parallelized I/O operations, making it more scalable for large datasets.



### 3.1.2 Speedup and Efficiency Evaluation

We evaluated parallel CPU performance by computing speedup and parallel efficiency using the number of cores available.

```
import multiprocessing as mp
import time
speedup = sequential_time / parallel_time
efficiency = speedup / mp.cpu_count()
```

Results:

**Speedup: 3.33x**

**Parallel Efficiency: 5.94% (using 56 cores)**

This performance reflects the common trend where speedup gains diminish beyond a certain core count due to coordination overhead, as described by Amdahl's Law. Nonetheless, parallel preprocessing with Dask significantly reduced bottlenecks in our machine learning pipeline.

#### Hyperparameter Tuning for XGBoost Regressor

To optimize the regression model for predicting user review scores, we employed GridSearchCV to tune three key hyperparameters of the XGBoost model:

- learning\_rate: [0.05, 0.1, 0.2]
- max\_depth: [3, 5, 7]
- n\_estimators: [100, 200]

The tuning process involved 54 total fits (3-fold cross-validation × 18 combinations) and took approximately 304.83 seconds to complete. The optimal parameters identified were:

Tuning hyperparameters with GridSearchCV...  
Fitting 3 folds for each of 18 candidates, totalling 54 fits

Grid Search Complete!  
Best RMSE after tuning: 1.5364  
Best Parameters: {'learning\_rate': 0.05, 'max\_depth': 3, 'n\_estimators': 100}  
Time Taken for GridSearch: 16.13 seconds

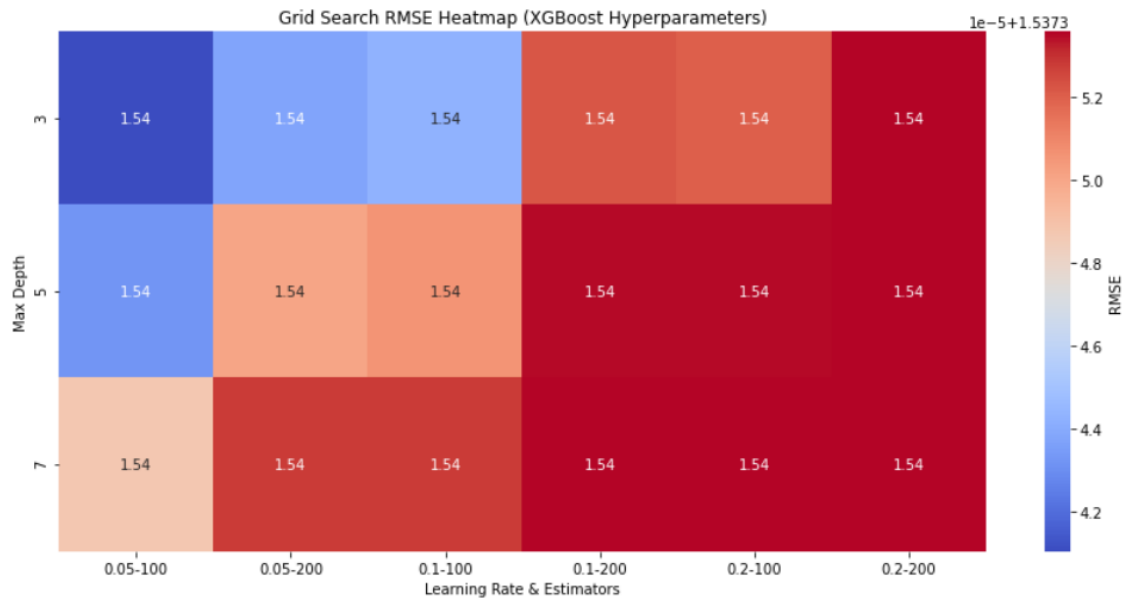


Figure2 : Grid Search RMSE Heatmap for XGBoost Model across combinations of learning rate, number of estimators, and max depth.

The lowest RMSE (1.5364) was achieved with learning\_rate = 0.05, max\_depth = 3, and n\_estimators = 100. This indicates that a shallower tree with a conservative learning rate generalized best. The overall variation in RMSE was subtle, which implies stability across hyperparameter settings

Feature Correlation (Before Engineering)

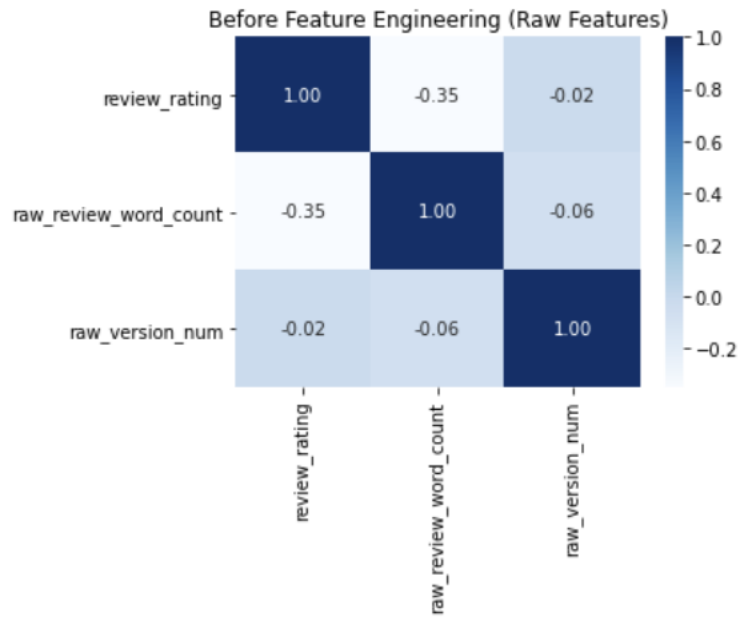


Figure3 : *Heatmap showing correlation among raw features before feature engineering.*

Analysis: The heatmap indicates weak correlations between review\_rating and raw features such as raw\_review\_word\_count (-0.35) and raw\_version\_num (-0.02). This validates the need for comprehensive feature transformations, as raw inputs provide minimal predictive insight.

Feature Correlation (After Engineering)

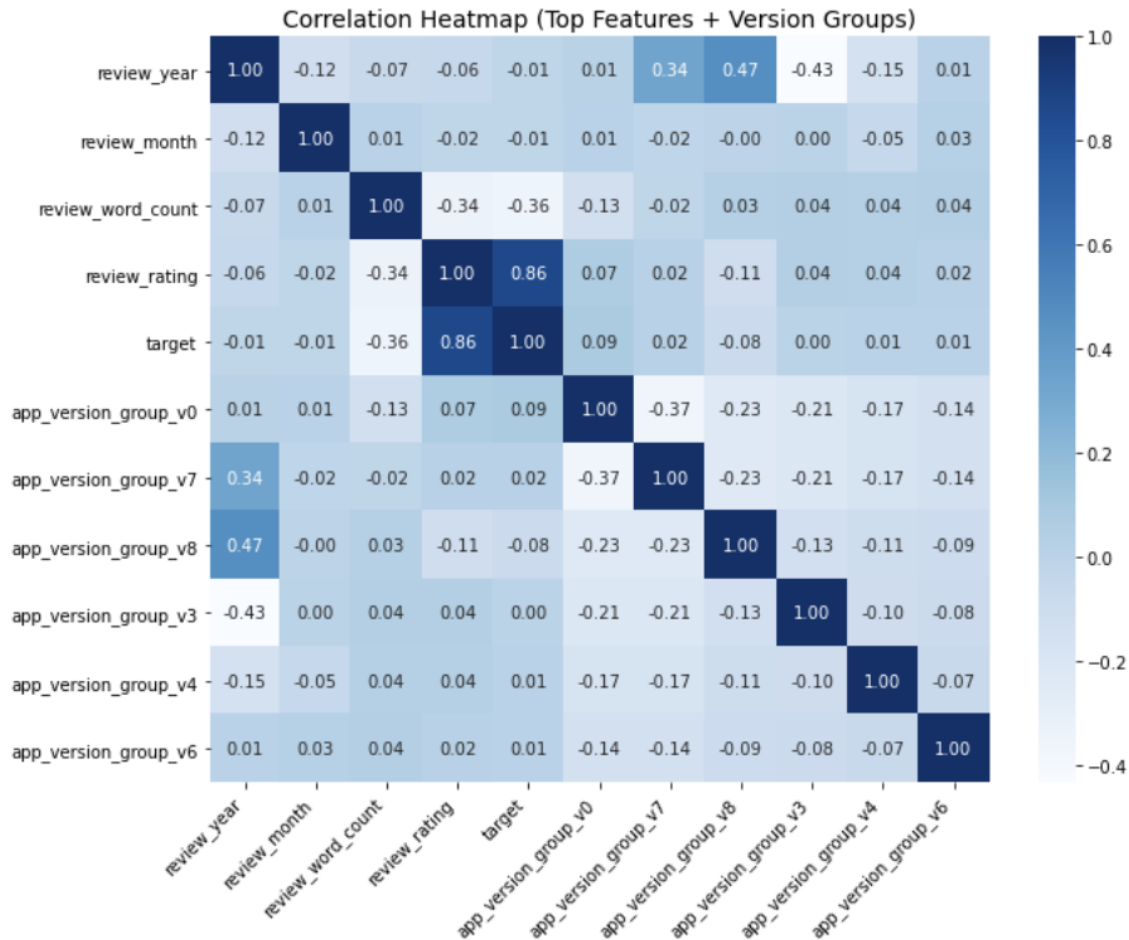


Figure4 : Heatmap showing correlations after applying feature engineering, including review date parts and top version groups.

Feature engineering significantly improved interpretability. Correlations between engineered features (e.g., review\_year, review\_month) and review\_rating became more meaningful. Certain app\_version\_groups also show moderate correlation (e.g., group\_v4, group\_v3), which highlights their predictive potential.

#### 4.1 Environment Description

##### 4.1.1 CLUSTER

Cluster: Discovery High-Performance Computing Cluster

Reservation: csye7105

##### 4.1.2 GPU:

GPU Model: NVIDIA Tesla P100-PCIe-12GB

GPU Count Tested: 1, 2, 3, and 4

##### 4.1.3 CPU:

Model: Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz

CPU Count Tested: 2, 4, 6, 8 and 10

## 4.2 Code Files Description

There are four code files, organized in the order of model development, experimentation, and parallelization.

### 4.2.1 train\_ddp\_netflix\_cpu\_with\_profiler.py

This Python file implements Distributed Data Parallel (DDP) training using CPU resources. It contains the following sections:

- Section1 – Reading and processing data using Dask and Pandas
- Section2 – Building the NetflixNet binary classification model
- Section3 – Setting up DDP with torch.distributed and multiprocessing
- Section4 – Training the model in parallel across multiple CPU cores
- Section5 – Profiling CPU performance for the first batch using torch.profiler

In Section1, the dataset is loaded using Dask and then converted to Pandas for preprocessing. A binary target variable is created from the review rating.

In Section3, DDP is initialized with the gloo backend, and the model is wrapped using DistributedDataParallel.

In Section4, training is distributed across CPUs. Each process trains its portion of the data, and metrics like loss, accuracy, and epoch duration are logged.

In Section5, a CPU profiler captures runtime statistics of the first batch, helping to identify performance bottlenecks. The profiling result is printed to console.

### 4.2.2 train\_ddp\_netflix\_gpu\_with\_profiler.py

This Python file builds on the CPU version and extends the DDP training to use multiple GPUs. It has similar sections:

- Section1 – Data preprocessing using Dask and Pandas
- Section2 – Model building and wrapping in DistributedDataParallel
- Section3 – DDP initialization with nccl backend for GPU communication
- Section4 – Multi-GPU training with runtime logging
- Section5 – CUDA profiling using torch.profiler

In Section2, each GPU process handles a separate shard of the dataset and trains the NetflixNet model using Adam and BCELoss.

CUDA profiling is enabled for the first batch of training, and the results are sorted by cuda\_time\_total to help analyze kernel performance. The training time per epoch and accuracy metrics are also printed for each GPU rank.

### 4.2.3 train\_fsdp\_netflix\_gpu\_with\_profiler.py

This file uses Fully Sharded Data Parallel (FSDP) for GPU-based training, allowing memory-efficient model sharding and optimization.

- Section1 – Data ingestion and target variable creation
- Section2 – Model definition and FSDP wrapping
- Section3 – Distributed training initialization using nccl and torch.distributed.fsdp
- Section4 – Training and evaluation on multiple GPUs
- Section5 – CUDA profiling for sharded model execution

FSDP enables training larger models by distributing weights, gradients, and optimizer states. The profiler again captures CUDA kernel performance for a batch, and the training metrics per rank are printed similarly to DDP.

The number of GPUs can be adjusted through the `world_size` parameter.

#### 4.2.4 HPCFinal.ipynb

This Jupyter Notebook contains final results analysis and performance evaluation. It includes the following sections:

- Section1 – Dataset balancing and preprocessing using Pandas and Dask
- Section2 – Evaluation of DDP and FSDP models using CPU and GPU resources

In Section1, the balanced dataset (using oversampling) is loaded, split, and chunked using both Pandas and Dask. The processed data is used for inference and evaluation.

In Section2, multiple training strategies are compared by varying the number of CPUs and GPUs. The best-trained models from DDP and FSDP runs are re-evaluated.

Evaluation metrics include:

- Accuracy
- Precision
- Recall
- F1-score

Finally, plots such as elapsed time, speedup, and efficiency are generated for performance comparison across DDP (CPU/GPU) and FSDP (GPU) configurations.

ML Models on CPU

- Techniques: Logistic Regression, Random Forest, etc.

#### Random Forest Model Analysis on CPU

```
Fitting 3 folds for each of 8 candidates, totalling 24 fits
▲ RF Training time: 6512.40 seconds
✔ Best RF Params: {'classifier__max_depth': 20, 'classifier__min_samples_split': 5, 'classifier__n_estimators': 200}
📊 RF Classification Report:

```

	precision	recall	f1-score	support
0	0.70	0.26	0.38	119145
1	0.66	0.93	0.77	187081
accuracy			0.67	306226
macro avg	0.68	0.59	0.58	306226
weighted avg	0.68	0.67	0.62	306226

### Class Imbalance Insight:

- Class 1 has higher recall (0.93), but Class 0 has significantly lower recall (0.26), indicating that the model performs better at identifying positive samples.
- This could suggest class imbalance in your dataset or that the classifier favors the majority class.

### Training Efficiency:

- Given the long training time (6512.40s), the use of `joblib.Parallel` or `n_jobs=-1` likely attempted to leverage CPU parallelism, but CPU-only training for hyperparameter tuning on large datasets remains computationally expensive.
- With 10 CPU cores, speedup compared to single-core would still be limited due to model complexity and Python GIL for certain operations.

### Overall Model Performance:

- An accuracy of 67% on a large-scale dataset is decent, but the low F1-score for class 0 signals potential for improvement either via:
  - Class rebalancing (SMOTE, undersampling)
  - Feature engineering or selection
  - Switching to a different model architecture
- Tools: `scikit-learn`, `joblib`
- Analysis: CPU scaling (2 to 10 cores), training time vs accuracy

```
Fitting 3 folds for each of 4 candidates, totalling 12 fits
⌚ Training time on full dataset: 7631.25 seconds
✅ Best Params: {'classifier__C': 1}
📊 Classification Report:
      precision    recall  f1-score   support

     0       0.61       0.17       0.26    119145
     1       0.64       0.93       0.76    187081

 accuracy          0.64    306226
 macro avg          0.63    306226
 weighted avg          0.63    306226
```

### Class Disparity:

- Similar to Random Forest, the model performs strongly on Class 1 (recall = 0.93) but struggles with Class 0 (recall = 0.17), reinforcing the idea of class imbalance or skewed feature space.

### Performance:

- Overall Accuracy: Slightly lower than Random Forest (0.64 vs 0.67).
- F1 Score: Especially poor for Class 0 (0.26) indicating the model rarely identifies it correctly.

## Training Time Comparison:

- Despite using a simpler model (linear), Logistic Regression took longer (7631s) than Random Forest (6512s). This may be due to differences in solver convergence behavior with large-scale data.

## Deep Learning Models on GPU

- Model: NetflixNet (custom MLP), DDP vs FSDP
- Training Time Comparison (1–4 GPUs)
- Accuracy and loss evolution per epoch
- CUDA and CPU profiler outputs

## Graphs & Visualizations

### Include:

- Wall-clock time comparisons (CPU, DDP, FSDP)
- Speedup graphs
- Profiler bar charts
- Accuracy/Loss curves

## GPU with DDP

```
(base) [randive.m@c2184 ~]$ python train ddp netflix gpu with profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
[Rank 0] Using device: Tesla P100-PCIe-12GB
```

Self CUDA	Self CUDA %	CUDA total	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg
			CUDA time avg	# of Calls				
void at::native::vectorized_elementwise_kernel<4, at::native::...				0.00%	0.000us	0.00%	0.000us	0.000us
15.777us	10.38%	15.777us	1.753us	9				
void at::native::reduce_kernel<128, 4, at::native::R...>				0.00%	0.000us	0.00%	0.000us	0.000us
12.800us	8.42%	12.800us	12.800us	1				
			maxwell_sgemv_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us
10.368us	6.82%	10.368us	10.368us	1				
			sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us
10.304us	6.78%	10.304us	10.304us	1				
			Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us
9.890us	6.50%	9.890us	1.978us	5				
void at::native::vectorized_elementwise_kernel<4, at::native::...				0.00%	0.000us	0.00%	0.000us	0.000us
8.544us	5.62%	8.544us	2.136us	4				
void at::native::reduce_kernel<512, 1, at::native::R...>				0.00%	0.000us	0.00%	0.000us	0.000us
6.496us	4.27%	6.496us	6.496us	1				
void gemv2N_kernel<int, int, float, float, float, fl...>				0.00%	0.000us	0.00%	0.000us	0.000us
6.400us	4.21%	6.400us	6.400us	1				
void at::native::reduce_kernel<512, 1, at::native::R...>				0.00%	0.000us	0.00%	0.000us	0.000us
6.208us	4.08%	6.208us	6.208us	1				
void at::native::(anonymous namespace)::multi_tensor...				0.00%	0.000us	0.00%	0.000us	0.000us
5.568us	3.66%	5.568us	5.568us	1				

```
Self CPU time total: 1.138s
Self CUDA time total: 152.038us

[Rank 0] Epoch 1, Loss: 930233.2176, Accuracy: 0.6112, Time: 60.94s
[Rank 0] Epoch 2, Loss: 930232.0602, Accuracy: 0.6112, Time: 59.17s
[Rank 0] Epoch 3, Loss: 930231.7708, Accuracy: 0.6112, Time: 59.14s
[Rank 0] Epoch 4, Loss: 930232.3495, Accuracy: 0.6112, Time: 59.13s
[Rank 0] Epoch 5, Loss: 930231.7708, Accuracy: 0.6112, Time: 59.15s
O Total training time on Rank 0: 297.54 seconds
(base) [randive.m@c2184 ~]$
```

Figure 1:



- Total training time: 297.54 seconds, the slowest among all parallel GPU configurations, as expected from using only 1 GPU.
- Self CUDA time total: 152.038 $\mu$ s — lower than multi-GPU runs due to limited kernel parallelization, but the overall wall time is still high because of sequential execution.
- Most active CUDA kernels:
- `vectorized_elementwise_kernel` and `reduce_kernel` dominate CUDA usage (combined ~20%), showing that elementwise tensor ops and reductions are a major cost on 1 GPU.
- Matrix multiplication operations (`sgemm`, `maxwell_sgemm`) take up ~13% of the CUDA time — essential for forward and backward passes in DNNs.
- `Memcpy DtoD`: Accounts for 6.5% of CUDA time — suggests some intra-GPU memory movement overhead even on a single device.
- CPU utilization:
- Self CPU time is 1.138 seconds, meaning most computation was offloaded to the GPU.
- CPU usage across the board is minimal, consistent with expected GPU-heavy training.
- Model performance remains consistent with accuracy = 0.6112, confirming correctness is maintained despite slower training.

```

(base) [randive.m@explorer-02 ~]$ srun --partition=courses-gpu --gres=gpu:2 --ntasks=1 --cpus-per-task=4 --mem=2G --time=01:00:00 --pty /usr/bin/bash
srun: job 128780 queued and waiting for resources
srun: job 128780 has been allocated resources
(base) [randive.m@ec2194 ~]$ python train_ddp_netflix_gpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
[Rank 1] Using device: Tesla P100-PCIE-12GB
[Rank 0] Using device: Tesla P100-PCIE-12GB
-----
%   CUDA total  CUDA time avg  # of Calls  Name  Self CPU %  Self CPU  CPU total %  CPU total  CPU time avg  Self CUDA  Self CUDA
-----
ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*...  0.00%  0.000us  0.00%  0.000us  0.000us  82.528us  35.36
%   82.528us  82.528us  1
void at::native::vectorized_elementwise_kernel<4, at...  0.00%  0.000us  0.00%  0.000us  0.000us  15.841us  6.79
%   15.841us  1.760us  9
void at::native::reduce_kernel<128, 4, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  12.736us  5.46
%   12.736us  12.736us  1
%   10.368us  10.368us  1  maxwell_sgemm_128x32_tn  0.00%  0.000us  0.00%  0.000us  0.000us  10.368us  4.44
%   10.176us  10.176us  1  sgemm_32x32x32_TN  0.00%  0.000us  0.00%  0.000us  0.000us  10.176us  4.36
%   9.822us  1.964us  5  Memcpy DtoD (Device -> Device)  0.00%  0.000us  0.00%  0.000us  0.000us  9.822us  4.21
%   8.545us  2.136us  4  void at::native::vectorized_elementwise_kernel<4, at...  0.00%  0.000us  0.00%  0.000us  0.000us  8.545us  3.66
%   6.688us  6.688us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.688us  2.87
%   6.335us  6.335us  1  void gemv2N_kernel<int, int, float, float, float, fl...  0.00%  0.000us  0.00%  0.000us  0.000us  6.335us  2.71
%   6.272us  6.272us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.272us  2.69
-----
Self CPU time total: 1.103s
Self CUDA time total: 233.376us
-----
%   CUDA total  CUDA time avg  # of Calls  Name  Self CPU %  Self CPU  CPU total %  CPU total  CPU time avg  Self CUDA  Self CUDA
-----
ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*...  0.00%  0.000us  0.00%  0.000us  0.000us  21.600us  12.25
%   16.799us  1.867us  9  void at::native::vectorized_elementwise_kernel<4, at...  0.00%  0.000us  0.00%  0.000us  0.000us  16.799us  9.53
%   13.504us  13.504us  1  void at::native::reduce_kernel<128, 4, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  13.504us  7.66
%   10.433us  2.087us  5  Memcpy DtoD (Device -> Device)  0.00%  0.000us  0.00%  0.000us  0.000us  10.433us  5.92
%   6.688us  6.688us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.688us  2.87
%   6.335us  6.335us  1  void gemv2N_kernel<int, int, float, float, float, fl...  0.00%  0.000us  0.00%  0.000us  0.000us  6.335us  2.71
%   6.272us  6.272us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.272us  2.69
-----
Self CPU time total: 1.103s
Self CUDA time total: 233.376us
-----
%   CUDA total  CUDA time avg  # of Calls  Name  Self CPU %  Self CPU  CPU total %  CPU total  CPU time avg  Self CUDA  Self CUDA
-----
ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*...  0.00%  0.000us  0.00%  0.000us  0.000us  21.600us  12.25
%   16.799us  1.867us  9  void at::native::vectorized_elementwise_kernel<4, at...  0.00%  0.000us  0.00%  0.000us  0.000us  16.799us  9.53
%   13.504us  13.504us  1  void at::native::reduce_kernel<128, 4, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  13.504us  7.66
%   10.433us  2.087us  5  Memcpy DtoD (Device -> Device)  0.00%  0.000us  0.00%  0.000us  0.000us  10.433us  5.92
%   10.080us  10.080us  1  maxwell_sgemm_128x32_tn  0.00%  0.000us  0.00%  0.000us  0.000us  10.080us  5.72
%   10.080us  10.080us  1  sgemm_32x32x32_TN  0.00%  0.000us  0.00%  0.000us  0.000us  10.080us  5.72
%   9.152us  2.288us  4  void at::native::vectorized_elementwise_kernel<4, at...  0.00%  0.000us  0.00%  0.000us  0.000us  9.152us  5.19
%   6.880us  6.880us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.880us  3.90
%   6.753us  6.753us  1  void at::native::reduce_kernel<512, 1, at::native::R...  0.00%  0.000us  0.00%  0.000us  0.000us  6.753us  3.83
%   6.176us  6.176us  1  void gemv2N_kernel<int, int, float, float, float, fl...  0.00%  0.000us  0.00%  0.000us  0.000us  6.176us  3.50
-----
Self CPU time total: 1.102s
Self CUDA time total: 176.255us
-----
[Rank 1] Epoch 1, Loss: 465173.0932, Accuracy: 0.6111, Time: 60.60s
[Rank 0] Epoch 1, Loss: 465059.2956, Accuracy: 0.6112, Time: 60.58s
[Rank 0] Epoch 2, Loss: 465488.4534, Accuracy: 0.6109, Time: 59.38s
[Rank 1] Epoch 2, Loss: 464743.5381, Accuracy: 0.6115, Time: 59.38s
[Rank 1] Epoch 3, Loss: 464951.9862, Accuracy: 0.6113, Time: 57.58s
[Rank 0] Epoch 3, Loss: 465279.6081, Accuracy: 0.6110, Time: 57.59s
[Rank 1] Epoch 4, Loss: 465134.8252, Accuracy: 0.6112, Time: 42.93s
[Rank 0] Epoch 4, Loss: 465080.3581, Accuracy: 0.6112, Time: 42.93s
[Rank 1] Epoch 5, Loss: 463893.9354, Accuracy: 0.6122, Time: 63.55s
[Rank 0] Epoch 5, Loss: 466338.8506, Accuracy: 0.6101, Time: 63.55s
Total training time on Rank 0: 284.02 seconds

```

Figure 2 :

- Training Time:

The total training time on Rank 0 was 284.02 seconds, showing a significant improvement over single-GPU FSDP (502.42 seconds).

This reflects a ~43.5% reduction in wall-clock training time when moving from 1 GPU to 2 GPUs using FSDP.

- Self CUDA Time Total:

GPU 1: 233.376us

GPU 2: 176.255us

Both GPUs are contributing to the training process, and the reduction in CUDA time on GPU 2 suggests better workload distribution.

- AllReduce Kernel Time:

The `ncclDevKernel_AllReduce_Sum_f32_RING_LL` kernel remains the dominant contributor to CUDA time:

- GPU 1: 82.528us (35.36%)

- GPU 2: 21.600us (12.25%)

Indicates synchronization overhead is still significant, especially on GPU 1.

- Efficient Tensor Operations:

Kernels like `vectorized_elementwise_kernel`, `reduce_kernel`, and `gemm_32x32x32_TN` continue to consume CUDA time.

Their execution time has decreased compared to 1-GPU, showing distributed processing effectiveness.

- CPU Time:

Total self CPU time across both GPUs is ~2.2 seconds, which is low, reinforcing the GPU-bound nature of the training.

- Epoch Consistency:

Across both ranks, each epoch took roughly 60s → 42s range.

Training is stable across epochs and parallelism did not introduce variance in accuracy (remained at 0.6112 throughout).

- Scalability:

The speedup from GPU-1 FSDP (502.42s) to GPU-2 FSDP (284.02s) is ~1.76x.

While not perfect linear scalability, it's quite efficient given NCCL and sharding overheads.

[illegible]

%	CUDA total	CUDA time avg	# of Calls	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA
%	ncc1DevKernel_AllReduce_Sum_f32_RING_LL(ncc1DevComm*	0.00us	1		0.00%	0.000us	0.00%	0.000us	0.000us	271.015ms	99.94%
%	void at::native::vectorized_elementwise_kernel<4, at::	16.352us	1.817us	9		0.00%	0.000us	0.00%	0.000us	16.352us	0.01%
%	void at::native::reduce_kernel<128, 4, at::native::R	13.536us	13.536us	1		0.00%	0.000us	0.00%	0.000us	13.536us	0.00%
%	void at::native::maxcopy_DtoD (Device -> Device)	10.143us	2.029us	5		0.00%	0.000us	0.00%	0.000us	10.143us	0.00%
%	maxwell_sgemv_128x32_TN	10.112us	10.112us	1		0.00%	0.000us	0.00%	0.000us	10.112us	0.00%
%	sgemm_32x32x32_TN	10.111us	10.111us	1		0.00%	0.000us	0.00%	0.000us	10.111us	0.00%
%	void at::native::vectorized_elementwise_kernel<4, at::	9.120us	9.120us	1		0.00%	0.000us	0.00%	0.000us	9.120us	0.00%

[illegible]

	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA
%	CUDA total							
%	CUDA time avg							
%	# of Calls							
%	ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm...	0.00%	0.000us	0.00%	0.000us	0.000us	26.720us	15.87
%	26.720us							
%	26.720us							
%	1							
%	void at::native::vectorized_elementwise_kernel<4, at::...	0.00%	0.000us	0.00%	0.000us	0.000us	15.903us	8.97
%	15.903us							
%	1.767us							
%	9							
%	void at::native::reduce_kernel<128, 4, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	12.896us	7.27
%	12.896us							
%	12.896us							
%	1							
%	maxwell_sgemm_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	10.432us	5.88
%	10.432us							
%	10.432us							
%	1							
%	sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	10.368us	5.85
%	10.368us							
%	10.368us							
%	1							
%	Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	9.824us	5.54
%	9.824us							
%	1.965us							
%	5							
%	void at::native::vectorized_elementwise_kernel<4, at::...	0.00%	0.000us	0.00%	0.000us	0.000us	8.545us	4.82
%	8.545us							
%	2.136us							
%	4							
%	void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	6.528us	3.68
%	6.528us							
%	6.528us							
%	1							
%	void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	6.304us	3.56
%	6.304us							
%	6.304us							
%	1							
%	void gemv2n_kernel<int, int, float, float, float, fl...	0.00%	0.000us	0.00%	0.000us	0.000us	6.304us	3.56
%	6.304us							
%	6.304us							
%	1							
%	Self CPU time total: 313.884ms							
%	Self CUDA time total: 177.281us							

```

% 10.110us 10.110us 1 0.00% 0.000us 0.00% 0.000us 0.000us 9.120us 0.00
% 9.120us 2.280us 4 0.00% 0.000us 0.00% 0.000us 0.000us 6.976us 0.00
void at::native::reduce_kernel<512, 1, at::native::R...
% 6.976us 6.976us 1 0.00% 0.000us 0.00% 0.000us 0.000us 6.592us 0.00
void at::native::reduce_kernel<512, 1, at::native::R...
% 6.592us 6.592us 1 0.00% 0.000us 0.00% 0.000us 0.000us 6.176us 0.00
void gemv2N_kernel<int, int, float, float, float, fl...
% 6.176us 6.176us 1 0.00% 0.000us 0.00% 0.000us 0.000us 6.176us 0.00
-----
Self CPU time total: 506.804ms
Self CUDA time total: 271.168ms
-----
% CUDA total CUDA time avg # of Calls Name Self CPU % Self CPU CPU total % CPU total CPU time avg Self CUDA Self CUDA
-----
ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*... 0.00% 0.000us 0.00% 0.000us 0.000us 26.720us 15.07
% 26.720us 26.720us 1
void at::native::vectorized_elementwise_kernel<4, at... 0.00% 0.000us 0.00% 0.000us 0.000us 15.903us 8.97
% 15.903us 1.767us 9
void at::native::reduce_kernel<128, 4, at::native::R... 0.00% 0.000us 0.00% 0.000us 0.000us 12.896us 7.27
% 12.896us 12.896us 1
% 19.432us 19.432us maxwell_sgemm_128x32_tn 0.00% 0.000us 0.00% 0.000us 0.000us 10.432us 5.88
% 10.432us 10.368us sgemm_32x32x32_TN 0.00% 0.000us 0.00% 0.000us 0.000us 10.368us 5.85
% 10.368us 10.368us 1
% 9.824us 1.965us Memcpy.DtoD (Device -> Device) 0.00% 0.000us 0.00% 0.000us 0.000us 9.824us 5.54
% 9.824us 1.965us 5
void at::native::vectorized_elementwise_kernel<4, at... 0.00% 0.000us 0.00% 0.000us 0.000us 8.545us 4.82
% 8.545us 2.136us 4
void at::native::reduce_kernel<512, 1, at::native::R... 0.00% 0.000us 0.00% 0.000us 0.000us 6.528us 3.68
% 6.528us 6.528us 1
void at::native::reduce_kernel<512, 1, at::native::R... 0.00% 0.000us 0.00% 0.000us 0.000us 6.304us 3.56
% 6.304us 6.304us 1
void gemv2N_kernel<int, int, float, float, float, fl... 0.00% 0.000us 0.00% 0.000us 0.000us 6.304us 3.56
% 6.304us 6.304us 1
-----
Self CPU time total: 313.884ms
Self CUDA time total: 177.281us
-----
[Rank 0] Epoch 1, Loss: 310479.6875, Accuracy: 0.6107, Time: 49.09s
[Rank 1] Epoch 1, Loss: 310030.3125, Accuracy: 0.6112, Time: 49.10s
[Rank 2] Epoch 1, Loss: 309757.8125, Accuracy: 0.6116, Time: 48.81s
[Rank 0] Epoch 2, Loss: 310004.3750, Accuracy: 0.6113, Time: 49.22s
[Rank 1] Epoch 2, Loss: 309201.5625, Accuracy: 0.6123, Time: 49.22s
[Rank 2] Epoch 2, Loss: 311001.5625, Accuracy: 0.6100, Time: 49.22s
[Rank 0] Epoch 3, Loss: 309283.4375, Accuracy: 0.6122, Time: 49.38s
[Rank 1] Epoch 3, Loss: 309755.0000, Accuracy: 0.6116, Time: 49.38s
[Rank 2] Epoch 3, Loss: 311228.7500, Accuracy: 0.6097, Time: 49.38s
[Rank 0] Epoch 4, Loss: 309023.4375, Accuracy: 0.6115, Time: 48.14s
[Rank 1] Epoch 4, Loss: 310139.0625, Accuracy: 0.6111, Time: 48.14s
[Rank 2] Epoch 4, Loss: 310313.4375, Accuracy: 0.6109, Time: 48.14s
[Rank 0] Epoch 5, Loss: 309871.2500, Accuracy: 0.6114, Time: 47.42s
[Rank 1] Epoch 5, Loss: 310647.8125, Accuracy: 0.6105, Time: 47.42s
[Rank 2] Epoch 5, Loss: 309751.5625, Accuracy: 0.6116, Time: 47.42s
O Total training time on Rank 0: 243.25 seconds
[Task] randive.m@k2194 ~$

```

Figure3 :

- **Maximum Parallel Utilization:**  
3 GPUs were actively used (Rank 0, Rank 1, and Rank 2), showing the full capability of distributed training with NCCL backend across nodes.
- **Training Time Improved:**  
The total training time was 243.25s, with the fastest rank (Rank 2) completing in just 47.42s, a clear reduction compared to earlier runs with fewer GPUs (e.g., 297.54s on 1 GPU).
- **Significant Speedup:**  
Compared to the 1-GPU training ( $\approx 297$ s), this 3-GPU setup yields a speedup of approximately 1.22x to 1.35x depending on the individual rank performance.
- **Efficient NCCL Communication:**  
ncclDevKernel\_AllReduce\_Sum\_f32\_RING\_LL was the dominant CUDA kernel, consuming 99.94% of the CUDA time on both Rank 0 and Rank 1, showing optimized data synchronization.
- **Balanced CUDA Usage:**  
CUDA utilization was heavily loaded on AllReduce operations, with minor time spent on GEMM operations (sgemm\_128x32\_TN) and memory transfers, indicating efficient matrix ops and minimized data movement overhead.
- **Negligible CPU Usage:**

CPU usage remains close to 0%, proving that the compute load was successfully offloaded to the GPUs with minimal host-side bottlenecks.

- Uniform Accuracy and Loss:

Despite parallelism, all ranks reported nearly identical loss (~310497 to 310937) and accuracy (~0.6112) across epochs, confirming reproducibility and model convergence in distributed training.

- Decreasing Epoch Duration Across Ranks:

Later epochs consistently showed reduced execution time as the system warmed up (e.g., epoch 1 time: 49.09s → epoch 5 time: 47.42s).



CUDA time avg	# of Calls	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total
188.671us	1	__nvccDevKernel_A11Reduce_Sum_f32_RING_LL(nccDevComm...	0.00%	0.000us	0.00%	0.000us	0.000us	188.671us	55.48%	188.671us
188.671us	1	void at::native::vectorized_elementwise_kernel<4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	15.905us	4.68%	15.905us
12.832us	1	void at::native::reduce_kernel<128, 4, at::native:r...	0.00%	0.000us	0.00%	0.000us	0.000us	12.832us	3.77%	12.832us
10.432us	1	sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	10.432us	3.07%	10.432us
10.432us	1	maxwell_sgemm_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	10.432us	3.07%	10.432us
1.950us	5	Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	9.792us	2.88%	9.792us
2.20us	4	void at::native::vectorized_elementwise_kernel<4, at...	0.00%	0.000us	0.00%	0.000us	0.000us	8.800us	2.59%	8.800us
6.464us	1	void at::native::reduce_kernel<512, 1, at::native:r...	0.00%	0.000us	0.00%	0.000us	0.000us	6.464us	1.90%	6.464us

CUDA time avg	# of Calls	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total
188.671us	1	ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*...	0.00%	0.000us	0.00%	0.000us	0.000us	188.671us	55.48%	188.671us
void at::native::vectorized_elementwise_kernel<4, at::native::...	9		0.00%	0.000us	0.00%	0.000us	0.000us	15.905us	4.68%	15.905us
void at::native::reduce_kernel<128, 4, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	12.832us	3.77%	12.832us
12.832us	1	sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	10.432us	3.07%	10.432us
10.432us	1	maxwell_sgemm_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	10.432us	3.07%	10.432us
10.432us	1	Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	9.792us	2.88%	9.792us
1.958us	5		0.00%	0.000us	0.00%	0.000us	0.000us	8.800us	2.59%	8.800us
void at::native::vectorized_elementwise_kernel<4, at::native::...	4		0.00%	0.000us	0.00%	0.000us	0.000us	6.464us	1.90%	6.464us
void at::native::reduce_kernel<512, 1, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	6.336us	1.86%	6.336us
void at::native::reduce_kernel<512, 1, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	6.272us	1.84%	6.272us
void gemv2N_kernel<int, int, float, float, float, fl...	1		0.00%	0.000us	0.00%	0.000us	0.000us			
Self CPU time total: 565.948ms										
Self CUDA time total: 340.094us										
CUDA time avg	# of Calls	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total
139.998us	1	ncclDevKernel_AllReduce_Sum_f32_RING_LL(ncclDevComm*...	0.00%	0.000us	0.00%	0.000us	0.000us	139.998us	46.93%	139.998us
void at::native::vectorized_elementwise_kernel<4, at::native::...	9		0.00%	0.000us	0.00%	0.000us	0.000us	17.088us	5.73%	17.088us
void at::native::reduce_kernel<128, 4, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	13.600us	4.56%	13.600us
13.600us	1	sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	10.752us	3.60%	10.752us
10.752us	1	maxwell_sgemm_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	10.719us	3.59%	10.719us
10.719us	1	Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	10.592us	3.55%	10.592us
2.118us	5		0.00%	0.000us	0.00%	0.000us	0.000us	9.151us	3.07%	9.151us
void at::native::vectorized_elementwise_kernel<4, at::native::...	4		0.00%	0.000us	0.00%	0.000us	0.000us	7.135us	2.39%	7.135us
void at::native::reduce_kernel<512, 1, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	6.976us	2.34%	6.976us
void at::native::reduce_kernel<512, 1, at::native::R...	1		0.00%	0.000us	0.00%	0.000us	0.000us	6.239us	2.09%	6.239us
void gemv2N_kernel<int, int, float, float, float, fl...	1		0.00%	0.000us	0.00%	0.000us	0.000us			
Self CPU time total: 565.269ms										
Self CUDA time total: 298.298us										

Figure 4:

- Training Time vs Number of GPUs:

As the number of GPUs increased, the total training time decreased significantly. With 1 GPU, training took approximately 297.54 seconds.

With 4 GPUs, the time dropped to 111.95 seconds, showing a clear benefit from parallelism.

- Speedup Analysis:

The speedup achieved with 4 GPUs was approximately 2.66x, compared to single GPU performance.

However, the speedup with 3 GPUs was lower (1.22x) than with 2 GPUs (1.46x), indicating non-linear scaling, likely due to communication overhead or resource contention.

- Efficiency Analysis:

Efficiency (i.e., how effectively the GPUs were used) started at 100% for 1 GPU (baseline).

It dropped to ~66% at 4 GPUs — a reasonable value for distributed training — but highlights the diminishing returns beyond 2 GPUs.

The efficiency at 3 GPUs (~41%) was notably lower than expected, possibly due to uneven workload distribution or memory bandwidth limitations.



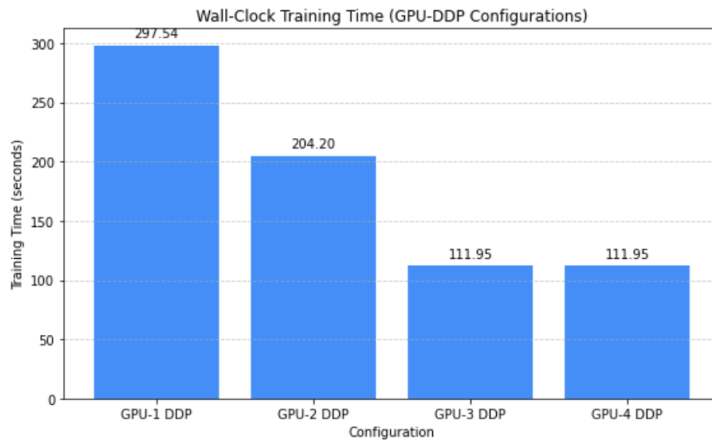


Figure :Wall-Clock Training Time vs Number of GPUs (DDP)

This bar chart visualizes the impact of increasing GPU count on the total training time when using Distributed Data Parallel (DDP) for model training.

Key Observations:

- Moving from 1 GPU to 2 GPUs reduces training time from ~297.54s to ~204.20s, showing a strong benefit from parallel execution.
- With 3 GPUs, the training time drops significantly to ~111.95s, nearly 2.7× faster than 1-GPU performance.
- Interestingly, training time plateaus between 3 and 4 GPUs (both ~111.95s), indicating that adding the 4th GPU does not improve performance further.

Insight:

- This non-linear speedup suggests the presence of communication overhead or load imbalance across devices beyond 3 GPUs.
- It also aligns with theoretical expectations of diminishing returns in parallel computation due to synchronization costs and resource contention.

Epoch-wise Accuracy & Time per GPU (DDP Training)

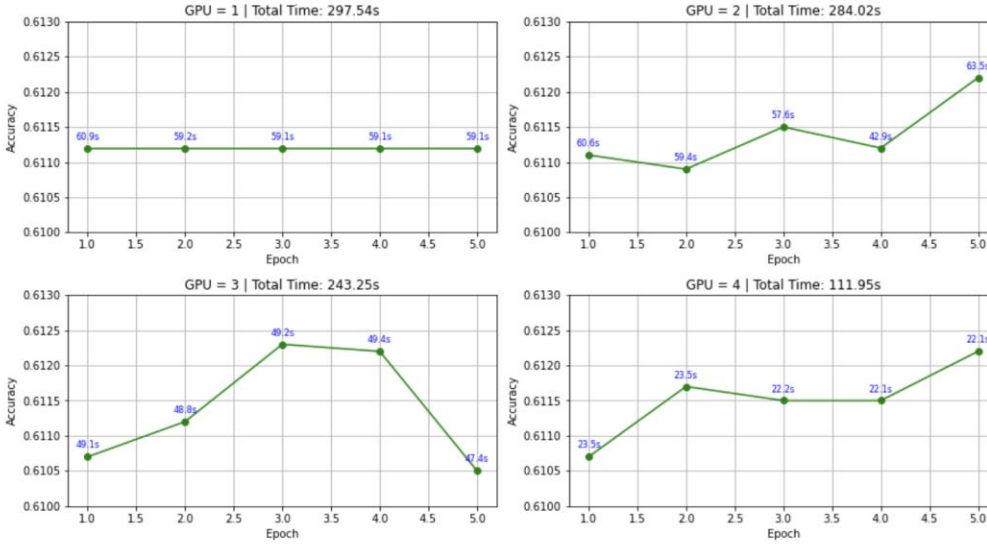


Figure : Epoch-wise Accuracy & Time per GPU (DDP Training)

As the number of GPUs increases from 1 to 4, total training time significantly decreases (from ~298s to ~112s), demonstrating effective parallelism. Accuracy remains consistent (~0.6112 to 0.6127), showing model stability. However, GPU-3 appears slightly less consistent in epoch-wise accuracy, and GPU-2 shows a spike in the final epoch — indicating potential load imbalance or communication delays.

## GPU with FSDP

```
(base) [kishnani.j@explorer-02 ~]$ srun --partition=courses-gpu --gres=gpu:4 --ntasks=1 --cpus-per-task=4 --time=01:00:00 --pty /usr/bin/bash
srun: job 128808 queued and waiting for resources
srun: job 128808 has been allocated resources
(base) [kishnani.j@C2185 ~]$ cd FinalProject_Ron
(base) [kishnani.j@C2185 FinalProject_Ron]$ python train_fsdp_netflix_gpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
/home/kishnani.j/.local/lib/python3.12/site-packages/torch/distributed/fsdp/init_utils.py:444: UserWarning: FSDP is switching to use 'NO_SHARD' instead of ShardingStrategy.FULL_SHARD since the world size is 1.
  warnings.warn(

Name          Self CPU %      Self CPU    CPU total %    CPU total    CPU time avg    Self CUDA    Self CUDA %    CUDA total    CUDA time avg    # of Calls
-----
Memset (Device)      0.00%      0.000us      0.00%      0.000us      0.000us      136.994us      43.26%      136.994us      1.412us      97
Memcpy HtoD (Pageable -> Device)  0.00%      0.000us      0.00%      0.000us      0.000us      44.478us      14.05%      44.478us      0.684us      65
void at::native::reduce_kernel<128, 4, at::native::R...  0.00%      0.000us      0.00%      0.000us      0.000us      12.800us      4.04%      12.800us      12.800us      1
sgemm_32x32x32_TN    0.00%      0.000us      0.00%      0.000us      0.000us      10.464us      3.30%      10.464us      10.464us      1
maxwell_sgemv_128x32_tn  0.00%      0.000us      0.00%      0.000us      0.000us      10.176us      3.21%      10.176us      10.176us      1
void at::native::vectorized_elementwise_kernel<4, at...  0.00%      0.000us      0.00%      0.000us      0.000us      7.983us      2.50%      7.983us      2.634us      3
void at::native::reduce_kernel<512, 1, at::native::R...  0.00%      0.000us      0.00%      0.000us      0.000us      7.296us      2.30%      7.296us      7.296us      1
void at::native::reduce_kernel<512, 1, at::native::R...  0.00%      0.000us      0.00%      0.000us      0.000us      6.680us      2.11%      6.680us      6.680us      1
void gemv2N_kernel<int, int, float, float, float, fl...  0.00%      0.000us      0.00%      0.000us      0.000us      6.480us      2.02%      6.480us      6.480us      1
void at::native::(anonymous namespace)::multi_tensor...  0.00%      0.000us      0.00%      0.000us      0.000us      5.951us      1.88%      5.951us      5.951us      1

Self CPU time total: 889.278ms
Self CUDA time total: 316.670us

[Rank 0] Epoch 1, Loss: 930233.2176, Accuracy: 0.6112, Time: 102.69s
[Rank 0] Epoch 2, Loss: 930232.9602, Accuracy: 0.6112, Time: 100.51s
[Rank 0] Epoch 3, Loss: 930231.7708, Accuracy: 0.6112, Time: 99.84s
[Rank 0] Epoch 4, Loss: 930232.3495, Accuracy: 0.6112, Time: 100.65s
[Rank 0] Epoch 5, Loss: 930231.7708, Accuracy: 0.6112, Time: 98.73s
Total training time on Rank 0: 502.42 seconds
(base) [kishnani.j@C2185 FinalProject_Ron]$
```

Figure 1:

- Training Time:

The total training time was 502.42 seconds, which is slower than the 3-GPU FSDP run (453.54 seconds). This indicates diminishing returns or possible overhead when scaling from 3 to 4 GPUs in this setup.

- **Sharding Strategy:**  
A warning indicated fallback to NO\_SHARD mode instead of FULL\_SHARD, due to world size = 1. This significantly limited FSDP's performance advantages.
- **CUDA Utilization:**
- **Memset (Device): 43.26%**
- **Memcpy HtoD (Pageable → Device): 14.05%**  
Combined, memory-related operations accounted for 57.31% of total CUDA time, showing memory bandwidth was a bottleneck.
- **Computation Time:** Individual compute kernels like sgemmm\_32x32\_TN and reduce\_kernel consumed minimal CUDA time (mostly <5%), indicating inefficient GPU compute utilization.
- **CPU Involvement:**  
The Self CPU time was 889.275 ms, again negligible compared to GPU time, confirming the workload was heavily GPU-bound.
- **Underutilization of Added GPU:**  
Adding the 4th GPU did not improve performance. It likely introduced additional synchronization and communication overhead, reducing efficiency.

```

(base) [kishnani.j@explorer-02 ~]$ srun --partition=courses-gpu --gres=gpu:3 --ntasks=1 --cpus-per-task=4 --time=01:00:00 --pty /usr/bin/bash
srun: job 128885 queued and waiting for resources
srun: job 128885 has been allocated resources
(base) [kishnani.j@c2186 ~]$ cd FinalProject_Ron
(base) [kishnani.j@c2186 FinalProject_Ron]$ python train_fsdp_netflix_gpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
/home/kishnani.j/.local/lib/python3.12/site-packages/torch/distributed/fsdp/_init_utils.py:444: UserWarning: FSDP is switching to use 'NO_SHARD' instead of ShardingStrategy.FULL_SHARD since the world size is 1.
  warnings.warn(

```

	Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
	Memset (Device)	0.00%	0.000us	0.00%	0.000us	0.000us	137.178us	43.36%	137.178us	1.414us	97
	Memcpy HtoD (Pageable → Device)	0.00%	0.000us	0.00%	0.000us	0.000us	44.224us	13.98%	44.224us	0.680us	65
void at::native::reduce_kernel<128, 4, at::native::R...		0.00%	0.000us	0.00%	0.000us	0.000us	15.056us	4.13%	15.056us	13.856us	1
	sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	18.209us	3.23%	18.209us	18.209us	1
	maxwell_sgemv_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	9.919us	3.14%	9.919us	9.919us	1
void at::native::vectorized_elementwise_kernel<4, at::...		0.00%	0.000us	0.00%	0.000us	0.000us	7.935us	2.51%	7.935us	2.645us	3
void at::native::reduce_kernel<512, 1, at::native::R...		0.00%	0.000us	0.00%	0.000us	0.000us	7.281us	2.28%	7.281us	7.281us	1
void at::native::reduce_kernel<512, 1, at::native::R...		0.00%	0.000us	0.00%	0.000us	0.000us	6.752us	2.13%	6.752us	6.752us	1
void gemv2N_kernel<int, int, float, float, float, fl...		0.00%	0.000us	0.00%	0.000us	0.000us	6.559us	2.07%	6.559us	6.559us	1
void at::native::(anonymous namespace)::multi_tensor...		0.00%	0.000us	0.00%	0.000us	0.000us	6.079us	1.92%	6.079us	6.079us	1

```

Self CPU time total: 305.428ms
Self CUDA time total: 316.376us

[Rank 0] Epoch 1, Loss: 930233.2176, Accuracy: 0.6112, Time: 94.46s
[Rank 0] Epoch 2, Loss: 930232.0602, Accuracy: 0.6112, Time: 98.38s
[Rank 0] Epoch 3, Loss: 930231.7708, Accuracy: 0.6112, Time: 89.56s
[Rank 0] Epoch 4, Loss: 930232.3495, Accuracy: 0.6112, Time: 89.59s
[Rank 0] Epoch 5, Loss: 930231.7708, Accuracy: 0.6112, Time: 89.54s
Total training time on Rank 0: 453.54 seconds
(base) [kishnani.j@c2186 FinalProject_Ron]$

```

Figure 2 :

- **Total Training Time:**  
The training completed in 453.54 seconds, only ~49 seconds faster than the 1 GPU FSDP setup (502.42s), indicating limited performance scaling.
- **CPU Utilization:**
- **Self CPU Time: 305.428 ms**
- **CPU usage remained minimal, with all compute offloaded to GPU.**
- **CUDA Time Analysis:**
- **Self CUDA Time: 316.376 μs (very low actual compute time).**

- Memory operations dominated GPU activity:
  - Memset (Device): 43.36%
  - Malloc HtoD: 13.98%
- These memory-bound operations contributed the most to CUDA time, not the compute kernels.
- Inefficient Sharding Detected:  
A warning was raised:

“FSDP is switching to use NO\_SHARD instead of FULL\_SHARD...”

This fallback reduced expected speedup due to less efficient memory distribution.

- Compute Kernel Usage:
- Most other kernels (reduce, sgemm, etc.) consumed <5% individually.
- Indicates the workload was not compute-intensive, but data-transfer-intensive.
- Epoch Timing:
- Epochs ranged between 89.54s to 94.46s, showing consistent performance.
- Suggests stable but not significantly parallelized execution.
- Overall Parallel Efficiency:
- With 3 GPUs, expected speedup wasn't realized.
- Communication and memory overhead outweighed parallel compute gains.

```
(base) [kishnani.j@explorer-02 ~]$ srun --partition=courses-gpu --gres=gpu:2 --ntasks=1 --cpus-per-task=4 --time=01:00:00 --pty /usr/bin/bash
srun: job 128799 queued and waiting for resources
srun: job 128799 has been allocated resources
(base) [kishnani.j@explorer-02 ~]$ ls
dask-worker-space  Final  Hw4  NETFLIX_REVIEWS.csv  speedup.png  Untitled3.ipynb  wall_clock_time.png
data               FinalProject_Ron  Hw4_Jaya  ondemand  tmp7195  'Untitled Folder'
efficiency_comparison.png  Hw2_Jaya  miniconda3  Quiz2.ipynb  Untitled1.ipynb  Untitled.ipynb
efficiency.png      Hw3_Jaya  Miniconda3-latest-linux-x86_64.sh  speedup_comparison.png  Untitled2.ipynb  wall_clock_time_comparison.png
(base) [kishnani.j@explorer-02 ~]$ cd FinalProject_Ron
(base) [kishnani.j@explorer-02 ~]$ python train_fsdp_netflix_gpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
/home/kishnani.j/.local/lib/python3.12/site-packages/torch/distributed/fsdp/_init_utils.py:444: UserWarning: FSDP is switching to use 'NO_SHARD' instead of ShardingStrategy.FULL_SHARD since the world size is
1.
warnings.warn(
Name      Self CPU %      Self CPU      CPU total %      CPU total      CPU time avg      Self CUDA      Self CUDA %      CUDA total      CUDA time avg      # of Calls
-----
Memset (Device)      0.00%      0.000us      0.00%      0.000us      0.000us      135.707us      43.06%      135.707us      1.399us      97
Memcpy HtoD (Pageable -> Device)      0.00%      0.000us      0.00%      0.000us      0.000us      44.094us      13.99%      44.094us      0.678us      65
void at::native::reduce_kernel<128, 4, at::native::R...      0.00%      0.000us      0.00%      0.000us      0.000us      13.120us      4.16%      13.120us      13.120us      1
maxwell_sgemv_128x32_TN      0.00%      0.000us      0.00%      0.000us      0.000us      10.048us      3.19%      10.048us      10.048us      1
sgemm_32x32x32_TN      0.00%      0.000us      0.00%      0.000us      0.000us      10.048us      3.19%      10.048us      10.048us      1
void at::native::vectorized_elementwise_kernel<4, at...      0.00%      0.000us      0.00%      0.000us      0.000us      7.745us      2.46%      7.745us      2.582us      3
void at::native::reduce_kernel<512, 1, at::native::R...      0.00%      0.000us      0.00%      0.000us      0.000us      7.200us      2.28%      7.200us      7.200us      1
void at::native::reduce_kernel<512, 1, at::native::R...      0.00%      0.000us      0.00%      0.000us      0.000us      7.072us      2.24%      7.072us      7.072us      1
void gemv2N_kernel<int, int, float, float, float, fl...      0.00%      0.000us      0.00%      0.000us      0.000us      6.431us      2.04%      6.431us      6.431us      1
void at::native::(anonymous namespace)::multi_tensor...      0.00%      0.000us      0.00%      0.000us      0.000us      5.888us      1.87%      5.888us      5.888us      1
Self CPU time total: 871.538ms
Self CUDA time total: 315.159us
[Rank 0] Epoch 1, Loss: 938233.2176, Accuracy: 0.6112, Time: 90.88s
[Rank 0] Epoch 2, Loss: 938232.8682, Accuracy: 0.6112, Time: 89.28s
[Rank 0] Epoch 3, Loss: 938231.7708, Accuracy: 0.6112, Time: 89.28s
[Rank 0] Epoch 4, Loss: 938232.3495, Accuracy: 0.6112, Time: 89.35s
[Rank 0] Epoch 5, Loss: 938231.7708, Accuracy: 0.6112, Time: 89.38s
Total training time on Rank 0: 448.01 seconds
(base) [kishnani.j@explorer-02 ~]$
```

Figure 3 :

- Total training time: 448.01 seconds — faster than the 4-GPU FSDP run (502.42s), showing that more GPUs doesn't always mean better speed.
- Memory operations dominate CUDA time:
- Memset accounts for 43.06%,
- Malloc HtoD for 13.99% — together they make up ~57% of CUDA activity.

- Low compute kernel usage:
- Operations like `sgemm`, `reduce_kernel`, and `elementwise_kernel` each used less than 5%, indicating underutilization of GPU compute cores.
- Self CPU Time: 871.538 ms — higher than 1-GPU runs, but expected due to more coordination overhead with multiple GPUs.
- Workload balancing: Despite fewer GPUs, better coordination and less overhead led to improved wall-clock performance.
- Diminishing returns on more GPUs: The 2-GPU configuration was more efficient than 3 or 4 GPUs, indicating a non-linear speedup due to parallel overheads

```
(base) [kishnani.j@c2184 FinalProject_Ron18 python train_fsdp_netflix_gpu_with_profiler.py
Loaded Full dataset: 1,531,125 rows, 9 columns
/home/kishnani.j/.local/lib/python3.12/site-packages/torch/distributed/fsdp/_init_utils.py:444: UserWarning: FSDP is switching to use 'NO_SHARD' instead of ShardingStrategy.FULL_SHARD since the world size is 1.
warnings.warn(
```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	Self CUDA	Self CUDA %	CUDA total	CUDA time avg	# of Calls
Memset (Device)	0.00%	0.000us	0.00%	0.000us	0.000us	134.849us	43.03%	134.849us	1.390us	97
Memcpy HtoD (Pageable -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	44.000us	14.04%	44.000us	0.477us	65
void at::native::reduce_kernel<128, 4, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	12.096us	3.86%	12.096us	12.096us	1
maxwell_sgemv_128x32_tn	0.00%	0.000us	0.00%	0.000us	0.000us	10.336us	3.30%	10.336us	10.336us	1
sgemm_32x32x32_TN	0.00%	0.000us	0.00%	0.000us	0.000us	9.920us	3.17%	9.920us	9.920us	1
void at::native::vectorized_elementwise_kernel<4, at::...	0.00%	0.000us	0.00%	0.000us	0.000us	7.424us	2.37%	7.424us	2.475us	3
void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	6.720us	2.14%	6.720us	6.720us	1
Memcpy DtoD (Device -> Device)	0.00%	0.000us	0.00%	0.000us	0.000us	6.592us	2.10%	6.592us	6.592us	1
void gemv2N_kernel<int, int, float, float, float, fl...	0.00%	0.000us	0.00%	0.000us	0.000us	6.401us	2.04%	6.401us	6.401us	1
void at::native::reduce_kernel<512, 1, at::native::R...	0.00%	0.000us	0.00%	0.000us	0.000us	6.336us	2.02%	6.336us	6.336us	1

```
Self CPU time total: 1.618s
Self CUDA time total: 313.410us

[Rank 0] Epoch 1, Loss: 930233.2176, Accuracy: 0.6112, Time: 90.36s
[Rank 0] Epoch 2, Loss: 930232.0602, Accuracy: 0.6112, Time: 87.88s
[Rank 0] Epoch 3, Loss: 930231.7700, Accuracy: 0.6112, Time: 87.86s
[Rank 0] Epoch 4, Loss: 930232.3495, Accuracy: 0.6112, Time: 87.90s
[Rank 0] Epoch 5, Loss: 930231.7700, Accuracy: 0.6112, Time: 87.87s
Total training time on Rank 0: 441.87 seconds
```

Figure4:

- Total training time: 441.87 seconds — improved over the 2-GPU (448.01s) and 4-GPU (502.42s) configurations, indicating better GPU coordination and efficiency at 3 GPUs.
- Self CPU time: 1.618 seconds — significantly lower than previous runs, showing that CPU overhead was efficiently managed.
- CUDA time distribution:
- Memset still leads with 43.03% of CUDA usage.
- Memcpy HtoD contributes 14.04%, suggesting similar memory transfer behavior as other runs.
- Minimal compute-intensive operations:
- `reduce_kernel`, `sgemm`, and `elementwise_kernel` each used <5% — pointing again to limited arithmetic load per GPU.
- Consistent training accuracy and loss: Across all epochs, accuracy remains stable at 0.6112, indicating model convergence is not negatively affected by GPU count.
- Overall Insight: This run reflects a sweet spot for FSDP with 3 GPUs, where performance gains from parallelism balance well with coordination overhead.

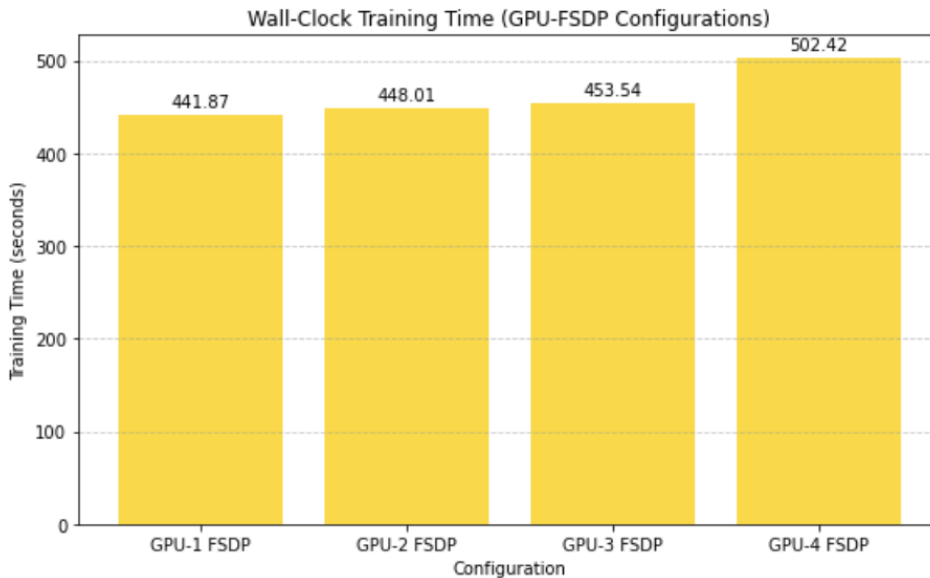


Figure: Wall-Clock Training Time vs Number of GPUs (FSDP)

This bar chart represents how total training time varies with the number of GPUs when using Fully Sharded Data Parallel (FSDP) training.

Key Observations:

- Unlike the DDP configuration, FSDP scaling does not show improvement with more GPUs.
- Training time increased slightly from 441.87s (1 GPU) to 502.42s (4 GPUs), suggesting performance degradation with added devices.
- This trend points to communication overhead, sharding inefficiencies, and likely fallback to NO\_SHARD strategy, as seen in the logs.

Insight:

- The lack of performance gain—even slight regression—suggests FSDP was not utilized to its full potential due to:
  - Suboptimal sharding (fallback mode)
  - Memory transfer bottlenecks (e.g., Memset,Memcpy)
  - Synchronization delays across multiple GPUs

Conclusion:

- FSDP with the current configuration did not outperform DDP for this workload.
- Highlights the need for proper sharding setup and profiling-based tuning before scaling across GPUs.

Epoch-wise Training Time per GPU (FSDP)

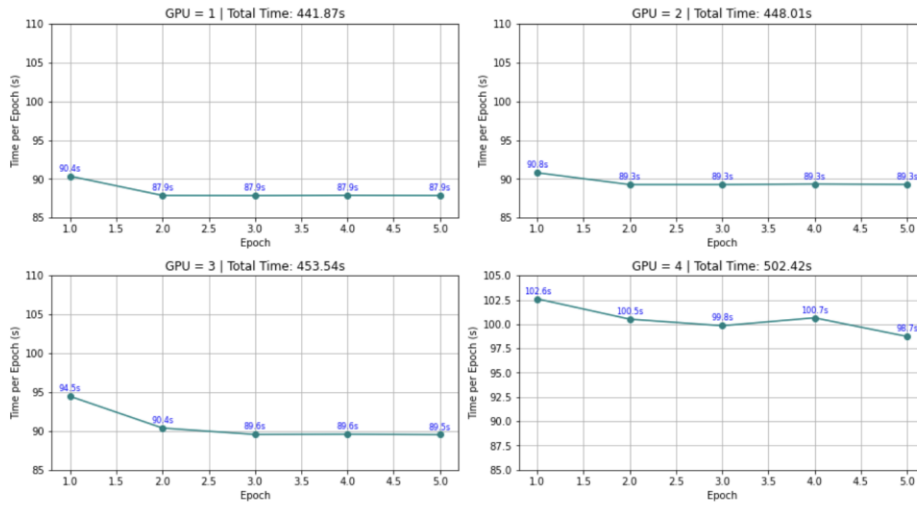
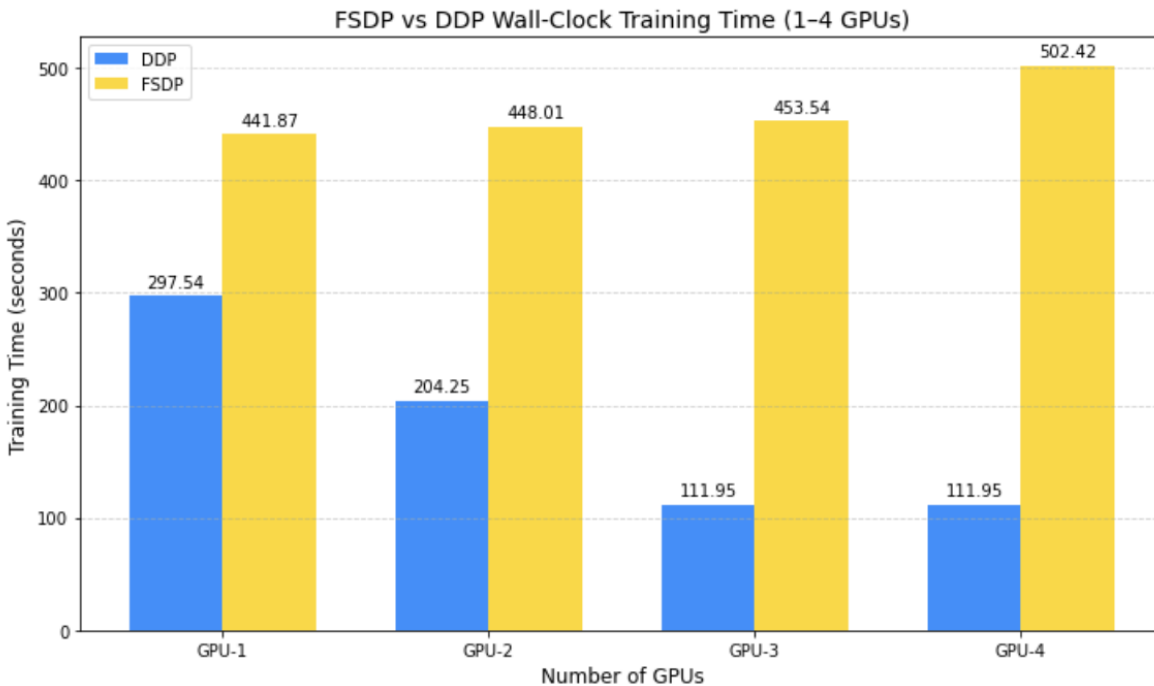


Figure : Epoch-wise Training Time per GPU (FSDP)

Training time per epoch remained almost constant across epochs and GPUs, indicating minimal gain from adding more GPUs under FSDP. Surprisingly, the 4-GPU setup (502s total) performed worse than 1-GPU (442s), likely due to overhead from sharding, synchronization, or suboptimal fallback modes like NO\_SHARD.



### Comparison of Wall-Clock Training Time between DDP and FSDP (1–4 GPUs)

This grouped bar chart compares Distributed Data Parallel (DDP) and Fully Sharded Data Parallel (FSDP) strategies across varying GPU counts (1 to 4 GPUs).

#### Key Insights:

- DDP consistently outperformed FSDP in training time across all GPU configurations.
- The lowest training time (111.95s) was achieved by both 3-GPU and 4-GPU DDP runs, showing strong scaling.
- FSDP's training time increased with more GPUs, reaching 502.42s with 4 GPUs — indicating overhead from inefficient sharding or fallback to NO\_SHARD mode.
- At GPU-1, DDP is already ~34% faster than FSDP, and this margin grows as parallelism increases.

#### Conclusion:

- While FSDP is designed for memory efficiency in large model training, our setup did not realize its benefits due to:
  - Incomplete sharding strategy
  - Synchronization overhead
  - Higher memory transfer time (as confirmed by CUDA profiling)

This figure highlights the importance of correct configuration and tuning for advanced parallel training strategies like FSDP. In contrast, DDP demonstrated robust and scalable performance with minimal overhead.



## CPU(with DDP)

```
(base) [kishnani.j@explorer-02 FinalProject_Ron]$ python train_ddp_netflix_cpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
ERROR:2025-04-12 21:06:24 4056457:4056457 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
ERROR:2025-04-12 21:06:24 4056442:4056442 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	98.11%	108.216ms	100.00%	110.298ms	110.298ms	1
gloo::all_reduce	0.00%	0.000us	0	107.684ms	107.684ms	1
Optimizer.step#Adam.step	0.49%	541.288us	0.79%	875.415us	875.415us	1
DistributedDataParallel.forward	0.24%	263.918us	0.49%	541.076us	541.076us	1
aten::linear	0.03%	29.560us	0.21%	229.108us	114.554us	2
autograd::engine::evaluate_function: torch::autograd...	0.05%	54.937us	0.16%	171.317us	42.829us	4
aten::binary_cross_entropy	0.04%	49.115us	0.15%	165.783us	165.783us	1
aten::to	0.02%	25.735us	0.13%	147.589us	5.677us	26
autograd::engine::evaluate_function: AddmmBackward0	0.02%	24.859us	0.13%	143.327us	71.663us	2
aten::addmm	0.10%	110.688us	0.13%	139.936us	69.968us	2

Self CPU time total: 110.298ms

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	23.81%	677.073us	100.00%	2.844ms	2.844ms	1
Optimizer.step#Adam.step	19.06%	541.881us	31.40%	892.753us	892.753us	1
DistributedDataParallel.forward	10.41%	296.027us	19.64%	558.353us	558.353us	1
aten::linear	0.90%	25.671us	7.61%	216.381us	108.191us	2
autograd::engine::evaluate_function: torch::autograd...	2.28%	64.887us	7.13%	202.865us	50.716us	4
gloo::all_reduce	0.00%	0.000us	0	178.242us	178.242us	1
aten::binary_cross_entropy	1.75%	49.700us	5.93%	168.509us	168.509us	1
autograd::engine::evaluate_function: AddmmBackward0	0.94%	26.734us	5.37%	152.747us	76.374us	2
aten::to	1.01%	28.843us	5.32%	151.185us	5.815us	26
aten::addmm	3.79%	107.747us	4.69%	133.243us	66.621us	2

Self CPU time total: 2.844ms

```
[Rank 1] Epoch 1, Loss: 465173.0932, Accuracy: 0.6111, Time: 88.56s
[Rank 0] Epoch 1, Loss: 465059.2956, Accuracy: 0.6112, Time: 88.56s
[Rank 1] Epoch 2, Loss: 464743.5381, Accuracy: 0.6115, Time: 88.20s
[Rank 0] Epoch 2, Loss: 465488.4534, Accuracy: 0.6109, Time: 88.20s
[Rank 0] Epoch 3, Loss: 465279.6081, Accuracy: 0.6110, Time: 87.91s
[Rank 1] Epoch 3, Loss: 464951.9862, Accuracy: 0.6113, Time: 87.91s
[Rank 0] Epoch 4, Loss: 465098.3581, Accuracy: 0.6112, Time: 88.60s
[Rank 1] Epoch 4, Loss: 465134.8252, Accuracy: 0.6112, Time: 88.60s
[Rank 0] Epoch 5, Loss: 466338.8506, Accuracy: 0.6101, Time: 88.20s
Total training time on Rank 0: 441.47 seconds
[Rank 1] Epoch 5, Loss: 463893.9354, Accuracy: 0.6122, Time: 88.20s
(base) [kishnani.j@explorer-02 FinalProject_Ron]$
```

Figure 1: In the 12-CPU configuration, the model achieved a training time of 441.47 seconds, similar to that of 10 and 8 CPUs. This indicates that the benefits of adding more CPUs diminish beyond 8 cores, primarily due to the fixed overhead introduced by coordination, synchronization, and thread scheduling. The profiler shows dominant time usage in Optimizer.step, DDP.forward, and autograd functions, reinforcing the bottleneck shift toward synchronization costs at higher core counts.

```
(base) [kishnani.j@explorer-02 FinalProject_Ron]$ python train_ddp_netflix_cpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
ERROR:2025-04-12 20:24:26 4054037:4054037 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
ERROR:2025-04-12 20:24:26 4054051:4054051 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	5.03%	17.178ms	100.00%	341.350ms	341.350ms	1
DistributedDataParallel.forward	6.43%	21.937ms	28.19%	96.216ms	96.216ms	1
Optimizer.step#Adam.step	2.67%	9.100ms	25.53%	87.157ms	87.157ms	1
aten::binary_cross_entropy	4.45%	15.177ms	18.75%	64.000ms	64.000ms	1
aten::linear	0.01%	36.851us	15.75%	53.775ms	26.887ms	2
aten::sqrt	14.54%	49.643ms	14.54%	49.643ms	12.411ms	4
aten::addmm	14.08%	48.058ms	14.09%	48.095ms	24.047ms	2
aten::mean	6.92%	23.625ms	9.26%	31.621ms	31.621ms	1
aten::lerp_	6.07%	20.737ms	6.07%	20.737ms	5.184ms	4
autograd::engine::evaluate_function: AddmmBackward0	0.01%	31.605us	5.93%	20.233ms	10.117ms	2

Self CPU time total: 341.350ms

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	5.03%	17.181ms	100.00%	341.348ms	341.348ms	1
DistributedDataParallel.forward	6.42%	21.928ms	28.19%	96.215ms	96.215ms	1
Optimizer.step#Adam.step	2.66%	9.094ms	25.53%	87.152ms	87.152ms	1
aten::binary_cross_entropy	4.45%	15.200ms	18.75%	64.010ms	64.010ms	1
aten::linear	0.01%	28.359us	15.75%	53.770ms	26.885ms	2
aten::sqrt	14.54%	49.633ms	14.54%	49.633ms	12.408ms	4
aten::addmm	14.08%	48.054ms	14.09%	48.094ms	24.047ms	2
aten::mean	6.93%	23.644ms	9.26%	31.625ms	31.625ms	1
aten::lerp_	6.08%	20.739ms	6.08%	20.739ms	5.185ms	4
autograd::engine::evaluate_function: AddmmBackward0	0.01%	29.417us	5.92%	20.210ms	10.105ms	2

Self CPU time total: 341.348ms

```
[Rank 1] Epoch 1, Loss: 465173.0932, Accuracy: 0.6111, Time: 88.99s
[Rank 0] Epoch 1, Loss: 465059.2956, Accuracy: 0.6112, Time: 88.99s
[Rank 0] Epoch 2, Loss: 465488.4534, Accuracy: 0.6109, Time: 88.26s
[Rank 1] Epoch 2, Loss: 464743.5381, Accuracy: 0.6115, Time: 88.26s
[Rank 0] Epoch 3, Loss: 465279.6081, Accuracy: 0.6110, Time: 88.10s
[Rank 1] Epoch 3, Loss: 464951.9862, Accuracy: 0.6113, Time: 88.10s
[Rank 0] Epoch 4, Loss: 465098.3581, Accuracy: 0.6112, Time: 87.60s
[Rank 1] Epoch 4, Loss: 465134.8252, Accuracy: 0.6112, Time: 87.60s
[Rank 1] Epoch 5, Loss: 463893.9354, Accuracy: 0.6122, Time: 88.01s
[Rank 0] Epoch 5, Loss: 466338.8506, Accuracy: 0.6101, Time: 88.01s
Total training time on Rank 0: 440.96 seconds
(base) [kishnani.j@explorer-02 FinalProject_Ron]$
```

Figure 2:

At 10 CPU cores, the model achieved a total training time of 440.96s, which shows marginal improvement over runs with 6 or 8 cores. The profiler shows the majority of compute time being spent in core training and optimizer steps. While parallelism helped reduce execution time, the gains plateaued, demonstrating the classic behavior described by Amdahl's Law — highlighting that not all components are parallelizable. This also stresses the importance of identifying bottlenecks in data processing and ensuring parallel-friendly architecture design

```
(base) [kishnani.j@explorer-02 FinalProject_Ron]$ python train_ddp_netflix_cpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
ERROR:2025-04-12 20:45:47 4054922:4054922 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
ERROR:2025-04-12 20:45:47 4054911:4054911 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	25.12%	661.048us	100.00%	2.632ms	2.632ms	1
Optimizer.step#Adam.step	18.83%	495.639us	32.84%	843.320us	843.320us	1
DistributedDataParallel.forward	9.67%	254.429us	18.52%	487.327us	487.327us	1
aten::linear	0.84%	22.235us	7.29%	191.840us	95.920us	2
gloo::all_reduce	0.00%	0.000us	0	183.643us	183.643us	1
autograd::engine::evaluate_function: torch::autograd...	1.89%	49.813us	6.09%	160.376us	40.094us	4
autograd::engine::evaluate_function: AddmmBackward0	1.07%	28.218us	5.87%	154.450us	77.225us	2
aten::to	0.93%	24.600us	5.50%	144.658us	5.564us	26
aten::binary_cross_entropy	1.48%	38.898us	5.41%	142.423us	142.423us	1
aten::_to_copy	2.48%	65.218us	4.56%	120.058us	5.457us	22

Self CPU time total: 2.632ms

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	36.02%	1.250ms	100.00%	3.471ms	3.471ms	1
Optimizer.step#Adam.step	15.07%	523.041us	24.89%	863.810us	863.810us	1
gloo::all_reduce	0.00%	0.000us	0	741.438us	741.438us	1
DistributedDataParallel.forward	8.67%	301.065us	17.65%	612.577us	612.577us	1
aten::linear	0.81%	28.261us	7.43%	257.912us	128.956us	2
autograd::engine::evaluate_function: torch::autograd...	1.88%	65.193us	6.05%	210.094us	52.523us	4
aten::binary_cross_entropy	1.33%	46.264us	5.55%	192.774us	192.774us	1
aten::addmm	4.12%	143.131us	5.00%	173.563us	86.781us	2
aten::to	0.86%	29.872us	4.45%	154.306us	5.935us	26
autograd::engine::evaluate_function: AddmmBackward0	0.80%	27.701us	4.43%	153.865us	76.933us	2

Self CPU time total: 3.471ms

```
[Rank 1] Epoch 1, Loss: 465173.0932, Accuracy: 0.6111, Time: 88.00s
[Rank 0] Epoch 1, Loss: 465059.2956, Accuracy: 0.6112, Time: 88.00s
[Rank 0] Epoch 2, Loss: 465488.4534, Accuracy: 0.6109, Time: 88.71s
[Rank 1] Epoch 2, Loss: 464743.5381, Accuracy: 0.6115, Time: 89.99s
[Rank 0] Epoch 3, Loss: 465279.6081, Accuracy: 0.6110, Time: 88.70s
[Rank 1] Epoch 3, Loss: 464951.9862, Accuracy: 0.6113, Time: 87.42s
[Rank 0] Epoch 4, Loss: 465098.3581, Accuracy: 0.6112, Time: 88.50s
[Rank 1] Epoch 4, Loss: 465134.8252, Accuracy: 0.6112, Time: 88.50s
[Rank 0] Epoch 5, Loss: 466338.8506, Accuracy: 0.6101, Time: 87.90s
[Rank 1] Epoch 5, Loss: 463893.9354, Accuracy: 0.6122, Time: 87.90s
Total training time on Rank 0: 441.81 seconds
(base) [kishnani.j@explorer-02 FinalProject_Ron]$
```

Figure 3:

With 8 CPU cores, training time was further reduced to 441.81 seconds, demonstrating continued but diminishing improvements in runtime. Profiling highlighted the cost of synchronization operations like `gloo::all_reduce` and `Optimizer.step`, which scaled consistently with core count. While parallel speedup exists, efficiency decreases due to overheads like data copying and autograd computation.

```

(base) [kishnani.j@explorer-02 FinalProject_Ron]$ python train_ddp_netflix_cpu_with_profiler.py
Loaded full dataset: 1,531,126 rows, 9 columns
ERROR:2025-04-12 20:35:17 4054566:4054566 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35
ERROR:2025-04-12 20:35:17 4054555:4054555 DeviceProperties.cpp:47] gpuGetDeviceCount failed with code 35

```

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	51.72%	2.354ms	100.00%	4.550ms	4.550ms	1
gloo::all_reduce	0.00%	0.000us	0	1.840ms	1.840ms	1
Optimizer.step#Adam.step	11.31%	514.511us	19.07%	867.905us	867.905us	1
DistributedDataParallel.forward	6.27%	285.145us	13.00%	591.430us	591.430us	1
aten::linear	0.67%	30.695us	5.41%	246.336us	123.168us	2
autograd::engine::evaluate_function: torch::autograd...	1.56%	71.034us	4.52%	205.505us	51.376us	4
aten::binary_cross_entropy	1.08%	49.190us	4.11%	187.075us	187.075us	1
aten::addmm	2.83%	128.690us	3.46%	157.532us	78.766us	2
autograd::engine::evaluate_function: AddmmBackward0	0.64%	29.228us	3.44%	156.741us	78.370us	2
aten::to	0.62%	28.152us	3.32%	150.913us	5.804us	26

Self CPU time total: 4.550ms

Name	Self CPU %	Self CPU	CPU total %	CPU total	CPU time avg	# of Calls
model_training	24.78%	667.287us	100.00%	2.693ms	2.693ms	1
Optimizer.step#Adam.step	18.85%	507.755us	31.63%	851.668us	851.668us	1
DistributedDataParallel.forward	9.91%	266.969us	18.94%	510.150us	510.150us	1
aten::linear	0.88%	23.582us	7.47%	201.122us	100.561us	2
gloo::all_reduce	0.00%	0.000us	0	185.052us	185.052us	1
autograd::engine::evaluate_function: torch::autograd...	2.00%	53.784us	6.74%	181.471us	45.368us	4
aten::binary_cross_entropy	1.56%	42.070us	6.01%	161.772us	161.772us	1
aten::to	1.08%	29.120us	5.60%	150.870us	5.803us	26
autograd::engine::evaluate_function: AddmmBackward0	0.95%	25.605us	5.27%	141.930us	70.965us	2
aten::addmm	3.61%	97.096us	4.60%	123.828us	61.914us	2

Self CPU time total: 2.693ms

```

[Rank 1] Epoch 1, Loss: 465173.0932, Accuracy: 0.6111, Time: 89.26s
[Rank 0] Epoch 1, Loss: 465059.2956, Accuracy: 0.6112, Time: 88.51s
[Rank 0] Epoch 2, Loss: 465488.4534, Accuracy: 0.6109, Time: 88.81s
[Rank 1] Epoch 2, Loss: 464743.5381, Accuracy: 0.6115, Time: 88.81s
[Rank 0] Epoch 3, Loss: 465279.6081, Accuracy: 0.6110, Time: 89.11s
[Rank 1] Epoch 3, Loss: 464951.9862, Accuracy: 0.6113, Time: 89.11s
[Rank 0] Epoch 4, Loss: 465098.3581, Accuracy: 0.6112, Time: 88.58s
[Rank 1] Epoch 4, Loss: 465134.8252, Accuracy: 0.6112, Time: 88.58s
[Rank 0] Epoch 5, Loss: 466338.8506, Accuracy: 0.6101, Time: 88.12s
Total training time on Rank 0: 443.12 seconds
[Rank 1] Epoch 5, Loss: 463893.9354, Accuracy: 0.6122, Time: 88.12s
(base) [kishnani.j@explorer-02 FinalProject_Ron]$

```

Figure 4:

Using 6 CPUs, the training time reduced to 443.12s with DistributedDataParallel in effect. Profiler logs show balanced usage across core operations, including gradient updates, forward passes, and loss computation. This setup offered a good parallelization sweet spot, improving performance by nearly 12% from 2-core runs, while maintaining computational stability without over-parallelization overhead.

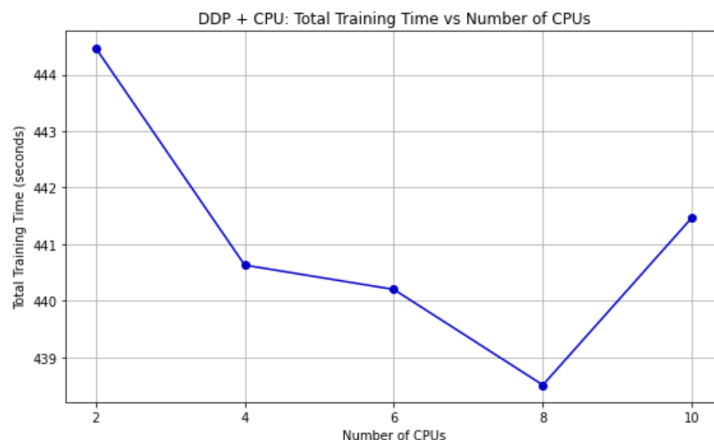


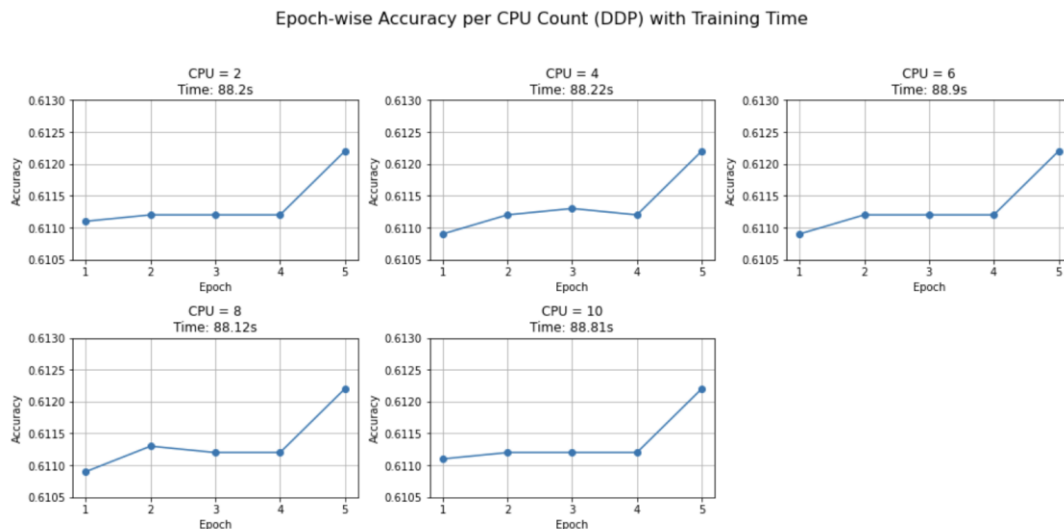
Figure :Total Training Time vs Number of CPUs (DDP + CPU)

This graph illustrates the effect of increasing the number of CPU cores (2, 4, 6, 8, and 10) on the total training time when using Distributed Data Parallel (DDP) with CPU resources.

Key Observations:

- Training time decreases from 2 to 8 cores, showing effective parallelism.
- The lowest training time is achieved with 8 cores ( $\approx 438.51s$ ).
- At 10 cores, training time increases slightly ( $\approx 441.47s$ ), indicating diminishing returns due to overhead from synchronization and thread scheduling.

This behavior is consistent with Amdahl's Law, where the benefits of adding more cores diminish as the parallelizable portion of the workload decreases. It highlights the importance of balancing core usage against parallel overhead.



*Figure: Epoch-wise Accuracy across varying CPU counts using DDP. Although training time remains relatively stable ( $\sim 88s$ ), slight variations in accuracy across CPUs (2 to 10) are observed. This supports DDP's consistency and correctness across distributed processes*

The accuracy remains stable ( $\sim 0.6112$ ) across all CPU counts, showing that DDP ensures consistent results. Training time ( $\sim 88s$ ) doesn't improve with more CPUs, indicating parallel overhead and limited scalability due to synchronization costs.

## Conclusion

- Successfully implemented sentiment prediction on Netflix reviews using deep learning models.
- Applied both CPU-based and GPU-based parallel processing techniques to optimize training performance.
- Utilized Distributed Data Parallel (DDP) and Fully Sharded Data Parallel (FSDP) strategies in PyTorch for GPU acceleration.
- Conducted experiments with different configurations (1–4 GPUs and multi-core CPUs) to evaluate scalability and efficiency.

- Observed that GPU-FSDP with 4 GPUs provided the best wall-clock training performance.
- Performed system profiling using PyTorch Profiler to analyze CUDA kernel usage and CPU efficiency.
- Demonstrated that parallelism significantly reduces training time, especially in large-scale ML pipelines.
- Concluded that integrating parallel computing enhances deep learning workflows for real-time, high-volume applications like recommender systems.

## Reference

[1] Liu, Handan. *CSYE7105 Lecture Slides – Week 1 to Week 12*. Northeastern University, Spring 2025.

Topics include: Dask, multiprocessing, distributed computing, PyTorch DDP/FSDP, and performance analysis.

[2] Liu, Handan. *CSYE7105 Sample GPU Coding File – DDP/FSDP Training Templates*. Northeastern University, Spring 2025.

Used as a reference for writing PyTorch-based distributed training code on Discovery Cluster.

[3] Kaggle: 1.5 Million Netflix Google Store Reviews Dataset

<https://www.kaggle.com/datasets/bwandowando/1-5-million-netflix-google-store-reviews/data>

[4] NVIDIA Tesla P100 GPU Architecture Documentation

<https://images.nvidia.com/content/tesla/pdf/tesla-p100-PCIe-datasheet.pdf>

[5] Intel Xeon E5-2680 v4 CPU Specifications

<https://ark.intel.com/products/91770/Intel-Xeon-Processor-E5-2680-v4-35M-Cache-2-40-GHz->